

EDISCO: EQUIVARIANT CONTINUOUS-TIME CATEGORICAL DIFFUSION FOR GEOMETRIC COMBINATORIAL OPTIMIZATION

000
001
002
003
004
005
006
007 **Anonymous authors**
008 Paper under double-blind review
009
010
011
012

ABSTRACT

013 Geometric combinatorial optimization problems, such as the Traveling Salesman
014 Problem (TSP), possess inherent symmetries under rotations, translations, and re-
015 flections in Euclidean space. These transformations are denoted as $E(2)$. How-
016 ever, existing neural network-based approaches, including recent diffusion-based
017 solvers, fail to exploit these geometric features. This paper presents EDISCO,
018 to the best of our knowledge, the first diffusion-based framework combining
019 $E(2)$ -equivariant graph neural networks with continuous-time categorical diffu-
020 sion models for solving geometric combinatorial problems. This approach intro-
021 duces an equivariant score network that respects geometric transformations while
022 operating on discrete edge variables, together with a continuous-time categorical
023 diffusion process that maintains $E(2)$ symmetries throughout the forward and re-
024 verse processes. By incorporating geometric awareness directly into the diffusion
025 process, EDISCO achieves notable improvements over the baseline. EDISCO re-
026 duces the state-of-the-art TSP optimality gaps on TSP-500 from 0.12% to 0.08%,
027 TSP-1000 from 0.30% to 0.22%, and TSP-10000 from 2.68% to 1.20%. EDISCO
028 demonstrates strong generalizability across problem sizes and also shows remark-
029 able efficiency, requiring only 33% to 50% of the training data compared to com-
030 peting diffusion methods across all problem scales.
031

1 INTRODUCTION

032 With diverse applications in logistics, circuit design, and resource allocation, geometric combi-
033 natorial optimization problems (GCOPs), such as the Traveling Salesman Problem (TSP), remain a
034 fundamental challenge in combinatorial optimization. Despite decades of research on exact and
035 heuristic solvers (Applegate et al., 2006; Helsgaun, 2017), the development of learning-based ap-
036 proaches has recently become the focus because of their potential for rapid inference and generaliza-
037 tion across various problem instances (Kool et al., 2019; Joshi et al., 2022; Fu et al., 2021). Recent
038 breakthroughs in diffusion models have opened new directions for solving GCOPs (Graikos et al.,
039 2022; Sun & Yang, 2023; Zhao et al., 2024). DIFUSCO (Sun & Yang, 2023) demonstrated graph-
040 based diffusion for TSP, while DISCO (Zhao et al., 2024) introduced residue-constrained generation
041 and analytical denoising to achieve up to 5.28× speedup over previous diffusion approaches.
042

043 However, an important observation is that TSP and similar GCOPs possess natural symmetries.
044 The solutions to these problems remain invariant under transformations in 2D Euclidean space, in-
045 cluding translations, rotations, and reflections (Ouyang et al., 2021; Bronstein et al., 2021). Such
046 transformations are denoted as $E(2)$. However, most existing neural network-based approaches fail
047 to capture the geometric structure of TSP. These methods require massive amounts of training data
048 and depend on high-quality optimal or near-optimal solutions for supervision, which are computa-
049 tionally expensive to obtain for large problems (Kool et al., 2019; Kwon et al., 2020; Joshi et al.,
050 2022). Furthermore, these models face significant memory and computational constraints. Even
051 moderate-scale problems cause memory overflow and require extensive training times (Bresson &
052 Laurent, 2021; Xin et al., 2021; Fu et al., 2021). This inefficiency originates from their need to learn
053 geometric invariances from scratch. Non-equivariant models attempt to address this issue with data
 augmentation, but this only shifts the problem. They require more training samples and still cannot
 ensure exact equivariance, especially on out-of-distribution data (Nordenfors et al., 2023; Esteves

et al., 2018). In contrast, models that explicitly incorporate geometric structure through equivariant architectures achieve better performance with less training data and smaller model sizes (Brehmer et al., 2024; Satorras et al., 2021; Batzner et al., 2022).

In addition to the geometric considerations, the choice of diffusion formulation presents another crucial design decision. While discrete diffusion models have shown promise for combinatorial problems (Sun & Yang, 2023; Austin et al., 2021), existing approaches employ discrete-time formulations with certain limitations. Although methods like DDIM (Song et al., 2021a) allow variable step counts at inference, they still rely on fixed discretization schemes that may accumulate approximation errors (Zhang & Chen, 2022b; Lu et al., 2022a; Ren et al., 2025). The discrete-time framework also limits access to adaptive numerical solvers that could dynamically adjust computational effort based on local dynamics (Ren et al., 2025; Zhang et al., 2024). Continuous-time formulations fundamentally address these limitations by treating the diffusion as a continuous process governed by stochastic differential equations (SDEs). This enables the use of numerical solvers with adaptive step sizes and higher-order integration methods that can achieve better accuracy with fewer number of function evaluations (NFEs) (Song et al., 2021a; Sun et al., 2023a).

In this work, we propose EDISCO, which leverages an Equivariant Graph Neural Network (EGNN) architecture to respect the geometric symmetries inherent in TSP instances. We also formulate the edge selection process as a continuous-time diffusion over discrete variables, enabling the derivation of analytical expressions for both the forward corruption process and the reverse denoising process, resulting in accelerated inference.

The novel contributions are as follows:

1. We introduce, to the best of our knowledge, the first continuous-time discrete diffusion model with built-in geometric equivariance for combinatorial optimization. Preserving problem symmetries improves both sample efficiency and solution quality.
2. We develop efficient training and sampling algorithms that leverage the analytical tractability of Continuous-time Markov Chains (CTMCs), compatible with higher-order accelerated solvers. These solvers achieve 2-3 \times speedups with better quality or up to 25 \times speedups for real-time applications compared to discrete-time methods.
3. The state-of-the-art performance is exceeded on TSP benchmarks (50-10000 cities). The optimality gap is reduced from 0.12% to 0.08% for TSP-500, from 0.30% to 0.22% for TSP-1000, and from 2.68% to 1.20% for TSP-10000. It only requires 33% to 50% of the training data compared to competing diffusion methods across all problem scales.
4. EDISCO is extended to solve real-world TSP problems and the Capacitated Vehicle Routing Problem (CVRP), outperforming SOTA approaches. These results validate the generalizability of EDISCO, indicating that it can be effectively applied to other GCOPs.

The remainder of this paper is organized as follows. Section 2 reviews related work on neural combinatorial optimization and diffusion models. Section 3 presents the continuous-time diffusion framework and the EGNN architecture. Section 4 presents comprehensive experimental results. Section 5 concludes with discussions and future directions.

2 RELATED WORK

2.1 NEURAL NETWORK-BASED TSP SOLVERS

Neural network-based approaches for TSP have recently become mainstream due to their potential for rapid inference, generalization across problem instances, and ability to learn from data without hand-crafted heuristics (Kool et al., 2019; Joshi et al., 2022; Fu et al., 2021). These approaches can be divided into autoregressive and non-autoregressive models. Autoregressive models (Kool et al., 2019; Kwon et al., 2020) construct solutions sequentially but require extensive training data and struggle with generalization (Joshi et al., 2022). Non-autoregressive approaches generate complete solutions simultaneously, evolving from limited heatmap representations (Joshi et al., 2019) to expressive diffusion models (Sun & Yang, 2023; Li et al., 2023; Yoon et al., 2024; Zhao et al., 2024).

108 Among diffusion approaches, DIFUSCO (Sun & Yang, 2023) pioneered graph-based diffusion for
 109 TSP, DISCO (Zhao et al., 2024) achieved 5.28 \times speedup through residue-constrained generation,
 110 and T2T (Li et al., 2023) improved quality via gradient-based search. CADO (Yoon et al., 2024)
 111 combines enhanced DIFUSCO by combining RL fine-tuning, but requires high-quality supervised
 112 data and expensive RL fine-tuning. The key insight is that all existing diffusion TSP solvers use
 113 discrete-time diffusion formulations and ignore TSP’s geometric structure.

114

115 2.2 GEOMETRIC DEEP LEARNING AND EQUIVARIANCE

116

117 Geometric deep learning leverages symmetries to improve efficiency and generalization (Bronstein
 118 et al., 2021). Equivariant neural networks ensure outputs transform consistently with symmetric in-
 119 puts, reducing sample complexity (Cohen & Welling, 2016). E(2)-equivariant models significantly
 120 improve TSP generalization (Ouyang et al., 2021), and geometric GNNs outperform standard archi-
 121 tectures for TSP tasks (Song et al., 2025). Theory confirms that equivariant models reduce training
 122 data requirements (Brehmer et al., 2024). Despite this evidence, most neural network-based TSP
 123 solvers lack geometric awareness. While Sym-NCO (Kim et al., 2022) uses regularizer-based sym-
 124 metry learning, it doesn’t achieve exact equivariance. These gaps motivate our approach as the first
 125 to combine exact E(2)-equivariance with diffusion.

126

127 2.3 CONTINUOUS-TIME DIFFUSION FORMULATIONS

128

129 Continuous-time formulations resolve fundamental limitations of discrete-time diffusion. CTMCs
 130 for discrete diffusion denoising (Campbell et al., 2022) enable analytical transition probabilities and
 131 flexible inference without retraining. Score-based continuous-time discrete diffusion (Sun et al.,
 132 2023b) provides better convergence properties and allows the use of higher-order integration meth-
 133 ods that significantly reduce the number of neural network evaluations (Song et al., 2021b). Though
 134 DiffUCO (Sankowski et al., 2024) applies continuous-time to unsupervised combinatorial opti-
 135 mization, it lacks geometric considerations and does not show results on large-scale geometric prob-
 136 lems. DISCO (Zhao et al., 2024) also achieves fast inference (1-2 steps) by replacing the entire
 137 numerical integration process with an analytically solvable form through decoupled diffusion mod-
 138 els (DDMs). This analytical solution completely bypasses the need for numerical ODE solvers but
 139 requires problem-specific residue constraints and sacrifices flexibility in the diffusion process.

140

141 3 METHOD

142

143 3.1 PROBLEM FORMULATION

144

145 GCOPs in Euclidean space are defined on a set of n nodes \mathcal{V} with coordinates $\{\mathbf{c}_i\}_{i=1}^n$, $\mathbf{c}_i \in \mathbb{R}^d$.
 146 The objective is to select a subset of edges or configurations, represented by a decision matrix
 147 $\mathbf{X} \in \{0, 1\}^{n \times n}$, that minimizes a distance-based cost function while satisfying problem-specific
 148 constraints. The objective is:

149

$$\mathbf{X}^* = \arg \min_{\mathbf{X}} f(\mathbf{X}, \{\mathbf{c}_i\}_{i=1}^n) \quad \text{s.t. } \mathbf{X} \in \mathcal{C} \quad (1)$$

150

where f is a distance-based cost function and \mathcal{C} represents the constraint.

151

In the case of TSP, given n cities with coordinates $\mathbf{c}_i \in \mathbb{R}^2$, we need to find a binary adjacency
 152 matrix $\mathbf{X} \in \{0, 1\}^{n \times n}$ where $X_{ij} = 1$ if edge (i, j) is included in the tour. The tour constraints
 153 are: each city has degree 2, and the selected edges form a connected cycle. We formulate this
 154 as a generative modeling problem Sun & Yang (2023); Li et al. (2023), learning the conditional
 155 distribution $p(\mathbf{X} | \{\mathbf{c}_i\}_{i=1}^n)$.

156

157

158 3.2 CONTINUOUS-TIME CATEGORICAL DIFFUSION FRAMEWORK

159

160

161

Unlike continuous diffusion models that operate in Euclidean space and require post-hoc quanti-
 162 zation, categorical diffusion directly models discrete decisions in their native space (Austin et al.,
 163 2021). This design choice eliminates quantization errors and ensures the model learns the true
 164 discrete distribution rather than a continuous approximation. DIFUSCO (Sun & Yang, 2023) also

162 demonstrated that categorical diffusion consistently outperforms continuous diffusion on TSP over
 163 all problem sizes.

164 Additionally, the continuous-time formulation offers extra benefits over discrete-time diffusion, as
 165 it enables exact likelihood computation, allows for flexible inference schedules without requiring
 166 retraining, and provides better theoretical guarantees for convergence (Campbell et al., 2022).
 167

168 **Forward Process** The forward process defines how clean data \mathbf{X}_0 progressively transitions to
 169 noise through a continuous-time Markov chain (CTMC). For K-state categorical variables, the in-
 170 stantaneous rate of transition between states is governed by the rate matrix (Campbell et al., 2022):
 171

$$\mathbf{Q}(t) = \beta(t) \left(\frac{1}{K} \mathbf{1} \mathbf{1}^T - \mathbf{I} \right) \quad (2)$$

172 where $\beta(t) = \beta_{\min} + t(\beta_{\max} - \beta_{\min})$ is a linear noise schedule with $t \in [0, 1]$. We set $\beta_{\min} = 0.1$
 173 and $\beta_{\max} = 1.5$.

174 The transition probability from time s to t is obtained by solving the Kolmogorov forward equa-
 175 tion (Norris, 1997), yielding the closed-form solution (Campbell et al., 2022):
 176

$$P_{ij}(t|s) = \frac{1}{K} + \left(\delta_{ij} - \frac{1}{K} \right) \exp \left(-K \int_s^t \beta(u) du \right) \quad (3)$$

177 For TSP with binary edge selection (K=2), this allows us to directly sample the noisy state \mathbf{X}_t from
 178 the clean data \mathbf{X}_0 :

$$P(\mathbf{X}_t = j | \mathbf{X}_0 = i) = P_{ij}(t|0) = \frac{1}{2} + \left(\delta_{ij} - \frac{1}{2} \right) \exp \left(-2 \int_0^t \beta(u) du \right) \quad (4)$$

179 For our linear schedule, the integral evaluates analytically to:
 180

$$\int_0^t \beta(u) du = \beta_{\min} t + \frac{1}{2} (\beta_{\max} - \beta_{\min}) t^2 \quad (5)$$

181 This closed-form expression enables exact sampling of \mathbf{X}_t given \mathbf{X}_0 at any time t without simulating
 182 intermediate states, crucial for efficient training. The exponential decay term ensures that as $t \rightarrow 1$,
 183 the transition probability approaches uniform ($P_{ij} \rightarrow 1/2$), completely corrupting the original signal
 184 while maintaining mathematical tractability.

185 **Reverse Process** The reverse process reconstructs clean data from noise by iteratively applying
 186 learned denoising steps. The key insight is that while the forward process is fixed and tractable,
 187 the reverse process requires learning the score function, which is the gradient of the log probability
 188 density. Using Bayes' rule, the posterior distribution for the reverse transition is:
 189

$$q(\mathbf{X}_{t-\Delta t} | \mathbf{X}_t, \mathbf{X}_0) = \frac{q(\mathbf{X}_t | \mathbf{X}_{t-\Delta t}, \mathbf{X}_0) q(\mathbf{X}_{t-\Delta t} | \mathbf{X}_0)}{q(\mathbf{X}_t | \mathbf{X}_0)} \quad (6)$$

200 Since the true \mathbf{X}_0 is unknown during inference, we need a neural network $s_\theta(\mathbf{X}_t, t, \{\mathbf{c}_i\})$ to predict
 201 it when given the noisy state and time. This parameterization, known as x_0 -prediction, is more
 202 stable than alternative parameterizations like noise prediction, especially in the low-noise region
 203 where reconstruction accuracy is significant (Salimans & Ho, 2022).

204 In EDISCO, we use an adaptive mixing strategy that dynamically balances between diffusion-based
 205 transitions and direct model predictions:

$$p_{\text{reverse}} = w(t) \cdot p_{\text{diffusion}} + (1 - w(t)) \cdot p_{\text{predicted}} \quad (7)$$

206 where $w(t) = t$ linearly decreases from 1 to 0 as the reverse process progresses. This design is
 207 motivated by the observation that early in the reverse process (large t), the noisy state contains little
 208 information about the target, making the diffusion dynamics essential for exploration. As t decreases
 209 and the signal emerges, direct predictions become increasingly reliable and should dominate to
 210 ensure precise reconstruction. For very small timesteps where $t < 0.1$ or $|\Delta t| < 0.02$, we switch
 211 entirely to deterministic transitions using the argmax of predicted probabilities. This strategy is
 212 evaluated in Appendix F.5.

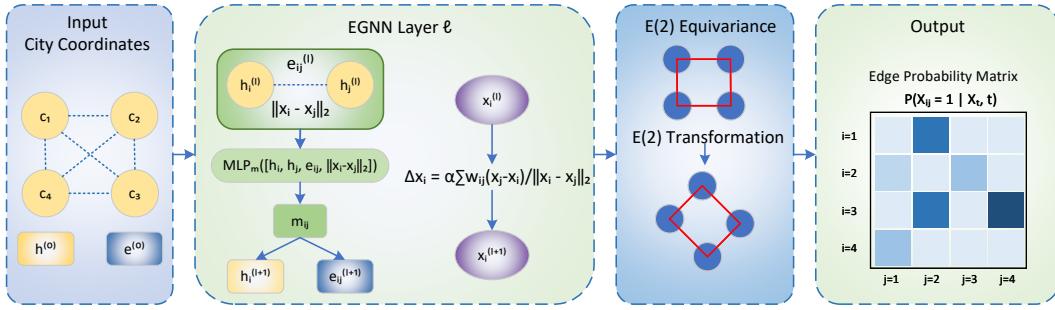


Figure 1: EDISCO’s EGNN Architecture Overview. EGNN layers process TSP instances while preserving $E(2)$ equivariance. The network outputs edge probabilities that remain invariant under geometric transformations.

3.3 EQUIVARIANT GRAPH NEURAL NETWORK ARCHITECTURE

Geometric Equivariance for TSP TSP possesses an inherent geometric structure that should be preserved. The $E(2)$ invariance of TSP solution has been recognized in prior work (Ouyang et al., 2021; Kim et al., 2022). If \mathbf{X}^* is optimal for cities $\{\mathbf{c}_i\}$, then \mathbf{X}^* remains optimal for transformed cities $\{g(\mathbf{c}_i)\}$ where $g \in E(2)$ represents any combinations of rotations, reflections, and translations. Traditional neural networks would need to learn this invariance from data, requiring extensive data augmentation and larger model capacity. However, we build equivariance directly into the architecture of EDISCO, ensuring that geometric transformations of inputs produce corresponding transformations of internal representations, so that the output can be maintained invariant (Thomas et al., 2018).

EGNN Layers with Stability Mechanisms We adapt the $E(n)$ -equivariant graph neural network (Satorras et al., 2021) with several crucial modifications for stable training on GCOPs. The architecture maintains three types of features: node features \mathbf{h}_i encoding local city information, edge features \mathbf{e}_{ij} representing pairwise relationships and tour decisions, and coordinate embeddings \mathbf{x}_i that evolve during message passing to capture geometric features. Figure 1 illustrates the EGNN architecture that maintains $E(2)$ -equivariance throughout the message passing process. The architecture processes TSP instances through multiple layers that preserve geometric symmetries while learning to predict edge probabilities for tour construction.

The message computation aggregates information from node pairs and their geometric relationship:

$$\mathbf{m}_{ij}^{(\ell)} = \text{MLP}_m \left([\mathbf{h}_i^{(\ell)}, \mathbf{h}_j^{(\ell)}, \mathbf{e}_{ij}^{(\ell)}, \|\mathbf{x}_i^{(\ell)} - \mathbf{x}_j^{(\ell)}\|_2] \right) \quad (8)$$

The inclusion of pairwise distances as scalar features, which are invariant under $E(2)$, allows the model to reason about geometric relationships without breaking equivariance.

Coordinate updates must preserve equivariance, achieved through the constrained form:

$$\Delta \mathbf{x}_i = \alpha \sum_{j \neq i} w_{ij} \cdot \frac{\mathbf{x}_j^{(\ell)} - \mathbf{x}_i^{(\ell)}}{\|\mathbf{x}_j^{(\ell)} - \mathbf{x}_i^{(\ell)}\|_2} \quad (9)$$

where the weights $w_{ij} = \tanh(\text{MLP}_c(\mathbf{m}_{ij}^{(\ell)})/\tau)$ control the influence of each neighbor. The temperature parameter $\tau = 10$ prevents saturation of the tanh function during early training when the MLP outputs may be large. The conservative step size $\alpha = 0.1$ is critical: larger values lead to coordinate collapse, where all cities converge to a single point. However, smaller values limit the model’s ability to learn useful geometric features. The normalization by distance ensures that the update magnitude is independent of the coordinate scale, therefore improving robustness. The selections of τ and α are extensively evaluated in Appendix F.6.

Edge features are updated with explicit time conditioning:

$$\mathbf{e}_{ij}^{(\ell+1)} = \text{LayerNorm}(\mathbf{e}_{ij}^{(\ell)} + \text{MLP}_e([\mathbf{e}_{ij}^{(\ell)}, \mathbf{m}_{ij}^{(\ell)}]) + \text{MLP}_t(\mathbf{t}_{\text{emb}})) \quad (10)$$

The time embedding \mathbf{t}_{emb} uses sinusoidal encoding (Vaswani et al., 2017). This enables the network to distinguish between fine-grained time differences near $t = 0$ and coarser differences at high noise levels.

Node features aggregate information from neighbors with gated attention:

$$\mathbf{h}_i^{(\ell+1)} = \text{LayerNorm}(\mathbf{h}_i^{(\ell)} + \text{MLP}_h([\mathbf{h}_i^{(\ell)}, \sum_{j \neq i} \sigma(\mathbf{m}_{ij}^{(\ell)}) \odot \mathbf{h}_j^{(\ell)}])) \quad (11)$$

Through the EGNN layers, EDISCO maintains exact E(2)-equivariance throughout the entire diffusion process. This is crucial for TSP because it reduces the effective complexity of the function to be learned.

Proposition 1. *Let $X = \mathbb{R}^{2n}$ denote the space of ordered 2D coordinates for n cities, and let $G = \text{E}(2)$ be the Euclidean transformation group on X by simultaneous rotation and translation of all city positions. Assume the transformation is free on the dataset (i.e., no non-trivial element $g \in G \setminus \{e\}$ fixes any configuration). Then:*

- (i) *The quotient space X/G is a smooth manifold of dimension $2n - 3$.*
- (ii) *Any G -equivariant function $F : X \rightarrow Y$ factors uniquely through the quotient as $F = \tilde{F} \circ \pi$, where $\pi : X \rightarrow X/G$ is the canonical projection and $\tilde{F} : X/G \rightarrow Y$ is a function on the quotient manifold.*
- (iii) *Learning a G -equivariant function is equivalent to learning a function on the $(2n - 3)$ -dimensional manifold X/G rather than on \mathbb{R}^{2n} .*

All the details about the notations can be found in the Default Notation. Although the dimension reduction from $2n$ to $2n - 3$ appears modest, its impact on learning is substantial. Equivariance forces the model to operate on the $(2n - 3)$ -dimensional orbit space instead of \mathbb{R}^{2n} , which reduces the metric entropy and the effective hypothesis-class complexity. The sample complexity reduction scales as $(1/\varepsilon)^3$ in covering number bounds, where ε is the desired approximation accuracy. We do not consider E(2) reflections because reflections are discrete transformations that do not further reduce the quotient dimension. For TSP, reflected tours are equivalent (same edges, only opposite tour sequence). The detailed proof of Proposition 1 is given in Appendix B.1. We also prove that the E(2)-Equivariance is preserved during the entire diffusion process in Appendix B.2.

3.4 INFERENCE AND TOUR DECODING

Solver Selection During inference, the continuous-time formulation allows for a flexible choice of accelerated and higher-order solvers without requiring retraining (Campbell et al., 2022). We extensively evaluate different solvers in Appendix F.1.

Tour Construction from Edge Probabilities The diffusion model outputs a probability matrix $P \in [0, 1]^{n \times n}$ where $P_{ij} = p(X_{ij} = 1)$ represents the model's confidence that edge (i, j) should be included in the optimal tour. Converting these soft probabilities to a valid discrete tour requires careful consideration of both model confidence and problem constraints.

Following the greedy decoding strategy from (Sun & Yang, 2023), we compute edge scores that balance model predictions with distance-based priors as $s_{ij} = (P_{ij} + P_{ji})/d_{ij}$, where the symmetrization $P_{ij} + P_{ji}$ accounts for TSP's undirected nature and d_{ij} denotes the Euclidean distance between nodes i and j .

The greedy construction algorithm maintains feasibility throughout the process. Starting with an empty tour, edges are processed in descending score order. An edge (i, j) is added if and only if both vertices have degree less than 2 (ensuring no city is visited more than twice) and adding the edge would not create a subtour (except for the final edge that completes the Hamiltonian cycle). Cycle detection is performed efficiently using a union-find data structure with path compression, achieving near-linear time complexity. The 2-opt local search (Lin & Kernighan, 1973) post-processing can be optionally applied to improve the tour.

324 **4 EXPERIMENTS**

325

326 **4.1 SETUP**

327

328 **Datasets** We follow the standard TSP evaluation protocol from Kool et al. (2019). Training instances are generated by sampling n cities uniformly from the unit square $[0, 1]^2$. We used the Concorde exact solver (Applegate et al., 2006) to generate datasets for TSP-50 and TSP-100, and used the LKH-3 (Helsgaun, 2017) heuristic solver for TSP-500 and TSP-1000. For evaluation, we use the standard test sets from Kool et al. (2019) for TSP-50/100 and Fu et al. (2021) for TSP-500 and above. EDISCO only requires 33% to 50% of the training data compared to baseline methods across all problem scales, which is reported in Appendix E.6.

334

335 **Graph Representation** For TSP-50/100, we use dense adjacency matrices representing complete graphs. For better scalability and fair comparisons, we apply graph sparsification to TSP-500 and above following the configuration of (Sun & Yang, 2023), with details displayed in Appendix E.4.

338

339 **Evaluation Metrics** We report three primary metrics: (1) average tour length, (2) average optimality gap, and (3) total run time. More experiment details can be found in Appendix E.1.

341

343 Table 1: Results on TSP-50 and TSP-100. RL: Reinforcement Learning, SL: Supervised Learning,
344 G: Greedy Decoding, 2O: 2-opt Post-processing. Concorde* represents the baseline for computing
345 the gap. All results except CADO are taken from Li et al. (2023). The results of CADO are taken
346 from Yoon et al. (2024).

| Algorithm | Type | TSP-50 | | TSP-100 | |
|--|-------------|---------------|-------------|----------------|-------------|
| | | Length ↓ | Gap ↓ | Length ↓ | Gap ↓ |
| Concorde* (Applegate et al., 2006) | Exact | 5.69 | 0.00 | 7.76 | 0.00 |
| 2-opt (Lin & Kernighan, 1973) | Heuristic | 5.86 | 2.95 | 8.03 | 3.54 |
| AM (Kool et al., 2019) | RL+G | 5.80 | 1.76 | 8.12 | 4.53 |
| GCN (Joshi et al., 2019) | SL+G | 5.87 | 3.10 | 8.41 | 8.38 |
| Transformer (Bresson & Laurent, 2021) | RL+G | 5.71 | 0.31 | 7.88 | 1.42 |
| POMO (Kwon et al., 2020) | RL+G | 5.73 | 0.64 | 7.87 | 1.07 |
| Sym-NCO (Kim et al., 2022) | RL+G | - | - | 7.84 | 0.94 |
| Image Diffusion (Graikos et al., 2022) | SL+G | 5.76 | 1.23 | 7.92 | 2.11 |
| DIFUSCO (Sun & Yang, 2023) | SL+G | 5.72 | 0.48 | 7.84 | 1.01 |
| T2T (Li et al., 2023) | SL+G | 5.69 | 0.04 | 7.77 | 0.18 |
| CADO (Yoon et al., 2024) | SL+RL+G | 5.69 | 0.01 | 7.77 | 0.08 |
| EDISCO with 50-step PNDM (ours) | SL+G | 5.69 | 0.01 | 7.76 | 0.04 |
| AM (Kool et al., 2019) | RL+G+2O | 5.77 | 1.41 | 8.02 | 3.32 |
| GCN (Joshi et al., 2019) | SL+G+2O | 5.70 | 0.12 | 7.81 | 0.62 |
| Transformer (Bresson & Laurent, 2021) | RL+G+2O | 5.70 | 0.16 | 7.85 | 1.19 |
| POMO (Kwon et al., 2020) | RL+G+2O | 5.73 | 0.63 | 7.82 | 0.82 |
| Sym-NCO (Kim et al., 2022) | RL+G+2O | - | - | 7.82 | 0.76 |
| DIFUSCO (Sun & Yang, 2023) | SL+G+2O | 5.69 | 0.09 | 7.78 | 0.22 |
| T2T (Li et al., 2023) | SL+G+2O | 5.69 | 0.02 | 7.76 | 0.06 |
| CADO (Yoon et al., 2024) | SL+RL+G+2O | 5.69 | 0.00 | 7.76 | 0.01 |
| EDISCO with 50-step PNDM (ours) | SL+G+2O | 5.69 | 0.00 | 7.76 | 0.01 |

371 **4.2 RESULTS**

372

373 For all EDISCO results shown in this section, we use the PNDM solver (Liu et al., 2022) with 50
374 steps, which achieves the best solution quality based on our extensive evaluation in Appendix F.1.

375

376 **TSP-50/100 Results** The results are shown in Table 1. EDISCO achieves near-optimal performance
377 with 0.01% gap on TSP-50 and 0.04% on TSP-100, substantially outperforming DIFUSCO
(0.48% and 1.01%) and T2T (0.04% and 0.18%). Unlike CADO, which requires both supervised

378
379
380
381
382

Table 2: Results on TSP-500 and TSP-1000. RL: Reinforcement Learning, SL: Supervised Learning, AS: Active Search, G: Greedy, S: Sampling, BS: Beam Search, 2O: 2-opt. Concorde* represents the baseline for computing the gap. All results except CADO and EDISCO are taken from Li et al. (2023). CADO results are from Yoon et al. (2024).

| Algorithm | Type | TSP-500 | | | TSP-1000 | | |
|---|------------|--------------|--------------|--------|--------------|--------------|--------|
| | | Length↓ | Gap↓ | Time | Length↓ | Gap↓ | Time |
| Concorde* (Applegate et al., 2006) | Exact | 16.55 | — | 37.66m | 23.12 | — | 6.65h |
| Gurobi (Gurobi Optimization, LLC, 2020) | Exact | 16.55 | 0.00% | 45.63h | — | — | — |
| LKH-3 (default) (Helsgaun, 2017) | Heuristics | 16.55 | 0.00% | 46.28m | 23.12 | 0.00% | 2.57h |
| AM (Kool et al., 2019) | RL+G | 20.02 | 20.99% | 1.51m | 31.15 | 34.75% | 3.18m |
| GCN (Joshi et al., 2019) | SL+G | 29.72 | 79.61% | 6.67m | 48.62 | 110.29% | 28.52m |
| POMO+EAS-Emb (Kwon et al., 2020) | RL+AS+G | 19.24 | 16.25% | 12.80h | — | — | — |
| POMO+EAS-Tab (Kwon et al., 2020) | RL+AS+G | 24.54 | 48.22% | 11.61h | 49.56 | 114.36% | 63.45h |
| DIMES (Qiu et al., 2022) | RL+G | 18.93 | 14.38% | 0.97m | 26.58 | 14.97% | 2.08m |
| DIMES (Qiu et al., 2022) | RL+AS+G | 17.81 | 7.61% | 2.10h | 24.91 | 7.74% | 4.49h |
| DIFUSCO (Sun & Yang, 2023) | SL+G | 18.11 | 9.41% | 5.70m | 25.72 | 11.24% | 17.33m |
| T2T (Li et al., 2023) | SL+G | 17.39 | 5.09% | 4.90m | 25.17 | 8.87% | 15.66m |
| CADO (Yoon et al., 2024) | SL+RL+G | 16.93 | 2.30% | 8.23m | 23.89 | 3.33% | 18.42m |
| EDISCO with 50-step PNNDM (ours) | SL+G | 16.87 | 1.95% | 2.19m | 23.78 | 2.85% | 6.84m |
| DIMES (Qiu et al., 2022) | RL+G+2O | 17.65 | 6.62% | 1.01m | 24.83 | 7.38% | 2.29m |
| DIMES (Qiu et al., 2022) | RL+AS+G+2O | 17.31 | 4.57% | 2.10h | 24.33 | 5.22% | 4.49h |
| DIFUSCO (Sun & Yang, 2023) | SL+G+2O | 16.81 | 1.55% | 5.75m | 23.55 | 1.86% | 17.52m |
| T2T (Li et al., 2023) | SL+G+2O | 16.68 | 0.78% | 4.98m | 23.41 | 1.25% | 15.90m |
| CADO (Yoon et al., 2024) | SL+RL+G+2O | 16.59 | 0.24% | 8.35m | 23.28 | 0.69% | 18.67m |
| EDISCO with 50-step PNNDM (ours) | SL+G+2O | 16.58 | 0.18% | 2.35m | 23.24 | 0.52% | 6.97m |
| EAN (Deudon et al., 2018) | RL+S+2O | 23.75 | 43.57% | 57.76m | 47.73 | 106.46% | 5.39h |
| AM (Kool et al., 2019) | RL+BS | 19.53 | 18.03% | 21.99m | 29.90 | 29.23% | 1.64h |
| GCN (Joshi et al., 2019) | SL+BS | 30.37 | 83.55% | 38.02m | 51.26 | 121.73% | 51.67m |
| DIMES (Qiu et al., 2022) | RL+S | 18.84 | 13.84% | 1.06m | 26.36 | 14.01% | 2.38m |
| DIMES (Qiu et al., 2022) | RL+AS+S | 17.80 | 7.55% | 2.11h | 24.89 | 7.70% | 4.53h |
| DIFUSCO (Sun & Yang, 2023) | SL+S | 17.48 | 5.65% | 19.02m | 25.11 | 8.61% | 59.18m |
| T2T (Li et al., 2023) | SL+S | 17.02 | 2.84% | 15.98m | 24.72 | 6.92% | 53.92m |
| CADO (Yoon et al., 2024) | SL+RL+S | 16.76 | 1.27% | 26.89m | 23.67 | 2.38% | 61.23m |
| EDISCO with 50-step PNNDM (ours) | SL+S | 16.72 | 1.05% | 7.82m | 23.57 | 1.95% | 23.27m |
| DIMES (Qiu et al., 2022) | RL+S+2O | 17.64 | 6.56% | 1.10m | 24.81 | 7.29% | 2.86m |
| DIMES (Qiu et al., 2022) | RL+AS+S+2O | 17.29 | 4.48% | 2.11h | 24.32 | 5.17% | 4.53h |
| DIFUSCO (Sun & Yang, 2023) | SL+S+2O | 16.69 | 0.83% | 19.05m | 23.42 | 1.30% | 59.53m |
| T2T (Li et al., 2023) | SL+S+2O | 16.61 | 0.37% | 16.03m | 23.30 | 0.78% | 54.67m |
| CADO (Yoon et al., 2024) | SL+RL+S+2O | 16.57 | 0.12% | 27.01m | 23.19 | 0.30% | 61.48m |
| EDISCO with 50-step PNNDM (ours) | SL+S+2O | 16.56 | 0.08% | 8.03m | 23.17 | 0.22% | 23.48m |

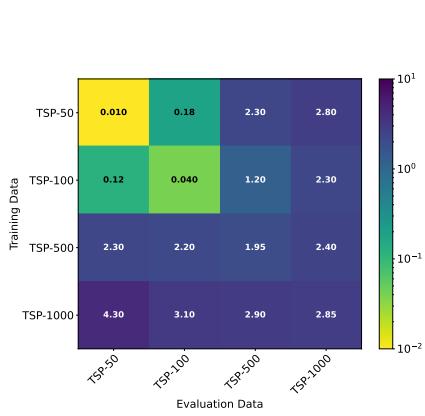
410
411
412413 and reinforcement learning, our purely supervised approach reaches comparable accuracy. With 2-
414 opt post-processing, EDISCO achieves optimal solutions (0.00% gap) on TSP-50 and 0.01% gap on
415 TSP-100, demonstrating that the equivariant architecture generates high-quality initial tours.429
430
431

Figure 2: Generalization performance across TSP sizes under greedy decoding.

TSP-500/1000 Results Table 2 presents results on larger-scale TSP instances. EDISCO achieves state-of-the-art performance across all decoding strategies. Using greedy decoding, EDISCO achieves gaps of 1.95% and 2.85% on TSP-500 and TSP-1000, outperforming DIFUSCO (9.41%, 11.24%) and T2T (5.09%, 8.87%). With 2-opt post-processing, EDISCO achieves near-optimal solutions with gaps of 0.18% and 0.52%, outperforming the previous best CADO (0.24%, 0.69%) while being 3.6x faster on TSP-500 and 2.7x faster on TSP-1000. The efficiency gain is particularly notable in sampling mode, where EDISCO requires only 7.82m and 23.27m compared to DIFUSCO’s 19.02m and 59.18m, demonstrating that the use of advanced numerical solvers significantly accelerates inference without compromising solution quality.

Generalization We study the generalization ability of EDISCO by training models on each problem scale from {TSP-50, TSP-100, TSP-500, TSP-1000} and evaluating them across all scales with only the greedy decoder. Figure 2 shows that EDISCO exhibits strong cross-size generalization, with models trained on TSP-1000 achieving gaps below 4.3% on all other problem scales, and particularly impressive performance of 2.90% on TSP-500. This generalizability outperforms other diffusion methods (Sun & Yang, 2023; Li et al., 2023).

Robustness to Training Data Variations

We evaluate EDISCO’s robustness to variations in training data quantity and quality, which are critical factors for practical deployment where obtaining optimal solutions may be computationally expensive. The left panel of Figure 3 illustrates that EDISCO maintains near-optimal performance even with limited data, achieving gaps below 0.07% with just 10% of training data, compared to 2.8% for DIFUSCO and 2.1% for T2T.

The right panel of Figure 3 examines model performance when trained on suboptimal solutions generated by the Farthest Insertion heuristic, which produces tours with an average gap of 7.5% to optimal on TSP-50 (Li et al., 2023). EDISCO achieves a 0.82% gap, outperforming DIFUSCO (2.75%) and T2T (1.35%). This experiment also only uses the greedy decoder for all methods.

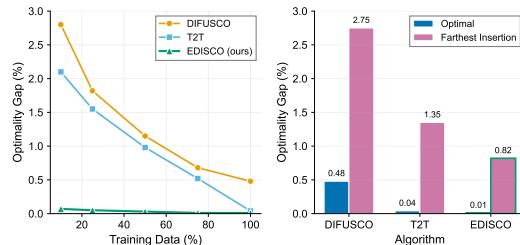


Figure 3: Results on TSP-50 performance. Left: Optimality gap as a function of training set size. Right: Performance comparison when trained on optimal data versus heuristic Farthest Insertion data.

Table 3: Ablation study on TSP-500 and TSP-1000. Each row removes one key component from the full EDISCO model.

| Model Variant | TSP-500 | | | TSP-1000 | | | Conv. Epoch |
|----------------------|--------------|-------------|--------------|--------------|-------------|--------------|-------------|
| | Length ↓ | Gap% ↓ | Time | Length ↓ | Gap% ↓ | Time | |
| EDISCO (Full) | 16.87 | 1.95 | 2.19m | 23.78 | 2.85 | 6.84m | 35 |
| w/o Equivariance | 17.14 | 3.58 | 2.31m | 24.29 | 5.06 | 7.15m | 48 |
| w/o Continuous-time | 17.02 | 2.86 | 4.43m | 24.11 | 4.28 | 15.58m | 42 |
| w/o Adaptive Mixing | 16.95 | 2.44 | 2.18m | 23.91 | 3.41 | 6.85m | 38 |

Ablation Studies Table 3 analyzes the contribution of each key component in EDISCO. Removing equivariance awareness causes the most significant degradation, increasing gaps from 1.95% to 3.58% on TSP-500 and from 2.85% to 5.06% on TSP-1000. It also requires 13 additional epochs to converge. The continuous-time formulation proves crucial for efficiency. Replacing it with DIFUSCO’s discrete-time approach doubles inference time (4.43m vs 2.19m on TSP-500) while degrading performance by 0.91% and 1.43% respectively. The adaptive mixing strategy shows the smallest but still meaningful impact, with its removal increasing gaps by 0.49% and 0.56%. These results confirm that all three components are significant, with equivariance providing the strongest inductive bias for learning geometric patterns, continuous-time enabling efficient sampling, and adaptive mixing ensuring accurate final reconstructions.

5 DISCUSSION AND CONCLUSION

We propose EDISCO, an equivariant continuous-time diffusion solver for GCOPs. By incorporating E(2) equivariance directly into the model architecture and formulating edge selection as continuous-time Markov chains, EDISCO learns geometric patterns more efficiently, enabling the use of advanced numerical solvers for fast inference. Future directions could include the exploration of adaptive step-size solvers and the theoretical analysis of convergence properties. Additionally, combining EDISCO with search-based refinement methods or integrating it with traditional optimization algorithms could further improve solution quality.

486 REPRODUCIBILITY STATEMENT
 487

488 We have made significant efforts to ensure reproducibility of our results. Details of the model and
 489 experimental settings are provided in the main text (Sections 3 and 4), as well as in the Appendix D
 490 and E. The source code and instructions for reproducing our experiments are available at <https://anonymous.4open.science/r/EDISCO>.
 491

492 REFERENCES
 493

494 Tara Akhound-Sadegh, Jarrid Rector-Brooks, Avishek Joey Bose, Sarthak Mittal, Pablo Castro,
 495 Yoshua Bengio, and Nikolay Malkin. Variational flow matching for discrete diffusion. In *International Conference on Machine Learning*, 2024.

496 David Applegate, Robert Bixby, Vasek Chvatal, and William Cook. Concorde TSP solver. Available
 497 at: <http://www.math.uwaterloo.ca/tsp/concorde>, 2006.
 500

501 Jacob Austin, Daniel D Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. Structured
 502 denoising diffusion models in discrete state-spaces. In *Advances in Neural Information
 503 Processing Systems*, volume 34, pp. 17981–17993, 2021.

504 Simon Batzner, Albert Musaelian, Lixin Sun, et al. E(3)-equivariant graph neural networks for
 505 data-efficient and accurate interatomic potentials. *Nature Communications*, 13(1):2453, 2022.

506 Irwan Bello, Hieu Pham, Quoc V Le, Mohammad Norouzi, and Samy Bengio. Neural combinatorial
 507 optimization with reinforcement learning. *arXiv preprint arXiv:1611.09940*, 2017.

508 Johann Brehmer, Sönke Behrends, Pim de Haan, and Taco Cohen. Does equivariance matter at
 509 scale? *arXiv preprint arXiv:2410.23179*, 2024.

510 Xavier Bresson and Thomas Laurent. The transformer network for the traveling salesman problem.
 511 *arXiv preprint arXiv:2103.03012*, 2021.

512 Michael M Bronstein, Joan Bruna, Taco Cohen, and Petar Veličković. Geometric deep learning:
 513 Grids, groups, graphs, geodesics, and gauges. *arXiv preprint arXiv:2104.13478*, 2021.

514 John Charles Butcher. *Numerical Methods for Ordinary Differential Equations*. John Wiley & Sons,
 515 3rd edition, 2016.

516 Andrew Campbell, Joe Benton, Valentin De Bortoli, Tom Rainforth, George Deligiannidis, and
 517 Arnaud Doucet. A continuous time framework for discrete denoising models. In *Advances in
 518 Neural Information Processing Systems*, 2022.

519 Andrew Campbell, Yuxin Zhang, Yaron Lipman, and George Deligiannidis. Discrete flow matching.
 520 *arXiv preprint arXiv:2407.12345*, 2024a.

521 Jonah Campbell, Yuxin Zhang, Yaron Lipman, and George Deligiannidis. Defog: Discrete flow
 522 matching for graph generation. *arXiv preprint arXiv:2410.04263*, 2024b.

523 Xinyun Chen and Yuandong Tian. Learning to perform local rewriting for combinatorial optimiza-
 524 tion. In *Advances in Neural Information Processing Systems*, volume 32, pp. 6281–6292, 2019.

525 Taco Cohen and Max Welling. Group equivariant convolutional networks. In *International Confer-
 526 ence on Machine Learning*, pp. 2990–2999, 2016.

527 Haoran Corsini, Zhigen Peng, and Brandon Amos. GLOP: Learning global partition and local
 528 construction for solving large-scale routing problems in real-time. In *Proceedings of the AAAI
 529 Conference on Artificial Intelligence*, volume 37, pp. 3740–3747, 2023.

530 Haoran Corsini, Zhigen Peng, and Brandon Amos. Glop: Learning global partition and local con-
 531 struction for solving large-scale routing problems in real-time. In *AAAI Conference on Artificial
 532 Intelligence*, 2024.

- 540 Michel Deudon, Pierre Cournut, Alexandre Lacoste, Yossiri Adulyasak, and Louis-Martin
 541 Rousseau. Learning heuristics for the tsp by policy gradient. In *International Conference on*
 542 *the Integration of Constraint Programming, Artificial Intelligence, and Operations Research*, pp.
 543 170–181. Springer, 2018.
- 544 Sander Dieleman, Laurent Sartran, Arman Roshannai, et al. Continuous diffusion for categorical
 545 data. *arXiv preprint arXiv:2211.15089*, 2022.
- 546 Darko Drakulic, Sofia Michel, Florian Mai, Arnaud Sors, and Jean-Marc Andreoli. Bq-nco: Bisim-
 547 ulation quotienting for generalizable neural combinatorial optimization. In *Advances in Neural*
 548 *Information Processing Systems*, volume 36, 2023.
- 549 Carlos Esteves, Christine Allen-Blanchette, Ameesh Makadia, and Kostas Daniilidis. Learning
 550 SO(3) equivariant representations with spherical CNNs. In *European Conference on Computer*
 551 *Vision*, pp. 52–68, 2018.
- 552 Zhang-Hua Fu, Kai-Bin Qiu, and Hongyuan Zha. Generalize a small pre-trained model to arbitrarily
 553 large TSP instances. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35,
 554 pp. 7474–7482, 2021.
- 555 Zhang-Hua Fu, Kai-Bin Qiu, and Hongyuan Zha. Neural combinatorial optimization with heavy
 556 decoder: Toward large scale generalization. *arXiv preprint arXiv:2310.07985*, 2024.
- 557 Maxime Gasse, Didier Chételat, Nicola Ferroni, Laurent Charlin, and Andrea Lodi. Exact combina-
 558 torial optimization with graph convolutional neural networks. In *Advances in Neural Information*
 559 *Processing Systems*, volume 32, 2019.
- 560 Yong Liang Goh, Zhiguang Cao, Yining Ma, Yanfei Dong, Mohammed Haroon Dupty, and Wee Sun
 561 Lee. Hierarchical neural constructive solver for real-world tsp scenarios. In *ACM SIGKDD*
 562 *Conference on Knowledge Discovery and Data Mining*, pp. 884–895, 2024.
- 563 Alexandros Graikos, Nikolay Malkin, Nebojsa Jojic, and Dimitris Samaras. Diffusion models as
 564 plug-and-play priors. In *Advances in Neural Information Processing Systems*, 2022.
- 565 Gurobi Optimization, LLC. Gurobi Optimizer Reference Manual, 2020. URL <https://www.gurobi.com>.
- 566 Keld Helsgaun. An extension of the Lin-Kernighan-Helsgaun TSP solver for constrained traveling
 567 salesman and vehicle routing problems. Technical report, Roskilde University, 2017.
- 568 Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *Advances*
 569 *in Neural Information Processing Systems*, volume 33, pp. 6840–6851, 2020.
- 570 Andre Hudson, Yong Liang Goh, Zhiguang Cao, and Wee Sun Lee. Self-improvement for neural
 571 combinatorial optimization: Sample without replacement, but improvement. *Transactions on*
 572 *Machine Learning Research*, 2024.
- 573 Chaitanya K Joshi, Thomas Laurent, and Xavier Bresson. An efficient graph convolutional network
 574 technique for the travelling salesman problem. *arXiv preprint arXiv:1906.01227*, 2019.
- 575 Chaitanya K Joshi, Quentin Cappart, Louis-Martin Rousseau, and Thomas Laurent. Learning the
 576 travelling salesperson problem requires rethinking generalization. *Constraints*, pp. 1–29, 2022.
- 577 Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-
 578 based generative models. In *Advances in Neural Information Processing Systems*, volume 35, pp.
 579 26565–26577, 2022.
- 580 Minsu Kim, Junyoung Park, and Jinkyoo Park. Sym-nco: Leveraging symmetricity for neural com-
 581 binatorial optimization. In *Advances in Neural Information Processing Systems*, volume 35, pp.
 582 1936–1949, 2022.
- 583 Wouter Kool, Herke Van Hoof, and Max Welling. Attention, learn to solve routing problems! In
 584 *International Conference on Learning Representations*, 2019.

- 594 Yeong-Dae Kwon, Jinho Choo, Byoungjip Kim, Iljoo Yoon, Youngjune Gwon, and Seungjai Min.
 595 POMO: Policy optimization with multiple optima for reinforcement learning. In *Advances in*
 596 *Neural Information Processing Systems*, volume 33, pp. 21188–21198, 2020.
- 597
 598 John M. Lee. *Introduction to Smooth Manifolds*, volume 218 of *Graduate Texts in Mathematics*.
 599 Springer, 2nd edition, 2012.
- 600 Yang Li, Jinpei Guo, Runzhong Wang, and Junchi Yan. T2T: From distribution learning in training
 601 to gradient search in testing for combinatorial optimization. In *Advances in Neural Information*
 602 *Processing Systems*, volume 36, 2023.
- 603 Shen Lin and Brian W Kernighan. An effective heuristic algorithm for the traveling-salesman prob-
 604 lem. *Operations Research*, 21(2):498–516, 1973.
- 605 Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching
 606 for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.
- 607 Luping Liu, Yi Ren, Zhijie Lin, and Zhou Zhao. Pseudo numerical methods for diffusion models on
 608 manifolds. In *International Conference on Learning Representations*, 2022.
- 609 Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, and Jun Zhu. Dpm-solver: A fast
 610 ode solver for diffusion probabilistic model sampling in around 10 steps. In *Advances in Neural*
 611 *Information Processing Systems*, volume 35, pp. 5775–5787, 2022a.
- 612 Cheng Lu, Yuhao Zhou, Fan Bao, Jianfei Chen, Chongxuan Li, and Jun Zhu. Dpm-solver++: Fast
 613 solver for guided sampling of diffusion probabilistic models. *arXiv preprint arXiv:2211.01095*,
 614 2022b.
- 615 Azalia Mirhoseini, Anna Goldie, Mustafa Yazgan, Joe Wenjie Jiang, Ebrahim Songhori, Shen Wang,
 616 Young-Joon Lee, Eric Johnson, Omkar Pathak, and Azadeh Nazi. A graph placement methodology
 617 for fast chip design. *Nature*, 594(7862):207–212, 2021.
- 618 Vinod Nair, Sergey Bartunov, Felix Gimeno, Ingrid von Glehn, Paweł Lichocki, Ivan Lobov, Bren-
 619 dan O’Donoghue, Nicolas Sonnerat, Christian Tjandraatmadja, and Pengming Wang. Solving
 620 mixed integer programs using neural networks. In *arXiv preprint arXiv:2012.13349*, 2020.
- 621 Mohammadreza Nazari, Afshin Oroojlooy, Lawrence Snyder, and Martin Takáć. Reinforcement
 622 learning for solving the vehicle routing problem. In *Advances in Neural Information Processing*
 623 *Systems*, volume 31, 2018.
- 624 Alexander Quinn Nichol and Prafulla Dhariwal. Improved denoising diffusion probabilistic models.
 625 In *International Conference on Machine Learning*, pp. 8162–8171, 2021.
- 626 Oskar Nordenfors, Fredrik Ohlsson, and Axel Flinth. Optimization dynamics of equivariant and
 627 augmented neural networks. *arXiv preprint arXiv:2303.13458*, 2023.
- 628 James R Norris. *Markov chains*. Cambridge University Press, 1997.
- 629 Wenbin Ouyang, Yisen Wang, Shaochen Han, et al. Generalization in deep RL for TSP problems
 630 via equivariance and local search. In *arXiv preprint arXiv:2110.03595*, 2021.
- 631 Ruizhong Qiu, Zhiqing Sun, and Yiming Yang. Dimes: A differentiable meta solver for combinato-
 632 rial optimization problems. In *Advances in Neural Information Processing Systems*, volume 35,
 633 pp. 25531–25546, 2022.
- 634 Gerhard Reinelt. Tsplib—a traveling salesman problem library. *ORSA Journal on Computing*, 3(4):
 635 376–384, 1991.
- 636 Yinuo Ren, Haoxuan Chen, Yuchen Zhu, Wei Guo, Yongxin Chen, Grant M Rotskoff, Molei Tao,
 637 and Lexing Ying. Fast solvers for discrete diffusion models: Theory and applications of high-
 638 order algorithms. *arXiv preprint arXiv:2502.00234*, 2025.
- 639 Tim Salimans and Jonathan Ho. Progressive distillation for fast sampling of diffusion models. *arXiv*
 640 *preprint arXiv:2202.00512*, 2022.

- 648 Sebastian Sanokowski, Sepp Hochreiter, and Sebastian Lehner. A diffusion model framework for
 649 unsupervised neural combinatorial optimization. In *International Conference on Machine Learn-*
 650 *ing*, 2024.
- 651
- 652 Simo Särkkä and Arno Solin. *Applied Stochastic Differential Equations*. Cambridge University
 653 Press, 2019.
- 654
- 655 Victor Garcia Satorras, Emiel Hoogeboom, and Max Welling. E(n) Equivariant Graph Neural Net-
 656 works. In *International Conference on Machine Learning*, pp. 9323–9332, 2021.
- 657
- 658 Kate Smith-Miles, Jano Van Hemert, and Xin Yu Lim. Understanding tsp difficulty by learning
 659 from evolved instances. *International Conference on Learning and Intelligent Optimization*, pp.
 266–280, 2010.
- 660
- 661 Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. In *Inter-
 662 national Conference on Learning Representations*, 2021a.
- 663
- 664 Ya Song, Laurens Bliek, and Yingqian Zhang. Geometrically invariant and equivariant graph neural
 networks for tsp algorithm selection and hardness prediction. 2025.
- 665
- 666 Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben
 667 Poole. Score-based generative modeling through stochastic differential equations. In *Inter-
 668 national Conference on Learning Representations*, 2021b.
- 669
- 670 Haoran Sun, Lijun Yu, Bo Dai, Dale Schuurmans, and Hanjun Dai. Score-based continuous-time
 discrete diffusion models. In *International Conference on Learning Representations*, 2023a.
- 671
- 672 Haoran Sun, Lijun Yu, Bo Dai, Dale Schuurmans, and Hanjun Dai. Score-based continuous-time
 discrete diffusion models. *International Conference on Learning Representations*, 2023b.
- 673
- 674 Zhiqing Sun and Yiming Yang. DIFUSCO: Graph-based diffusion solvers for combinatorial opti-
 675 mization. In *Advances in Neural Information Processing Systems*, 2023.
- 676
- 677 Nathaniel Thomas, Tess Smidt, Steven Kearnes, Lusann Yang, Li Li, Kai Kohlhoff, and Patrick
 678 Riley. Tensor field networks: Rotation-and translation-equivariant neural networks for 3d point
 679 clouds. 2018.
- 680
- 681 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez,
 Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Infor-
 682 mation Processing Systems*, pp. 5998–6008, 2017.
- 683
- 684 Oriol Vinyals, Meagan Fortunato, and Navdeep Jaitly. Pointer networks. In *Advances in Neural
 685 Information Processing Systems*, pp. 2692–2700, 2015.
- 686
- 687 Haoyu Peter Wang and Pan Li. Unsupervised learning for combinatorial optimization needs meta
 learning. In *International Conference on Learning Representations*, 2023.
- 688
- 689 Liang Xin, Wen Song, Zhiguang Cao, and Jie Zhang. Multi-decoder attention model with embedding
 690 glimpse for solving vehicle routing problems. *Proceedings of the AAAI Conference on Artificial
 691 Intelligence*, 35(13):12042–12049, 2021.
- 692
- 693 Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural
 networks? In *International Conference on Learning Representations*, 2019.
- 694
- 695 Keyulu Xu, Jingling Li, Mozhi Zhang, Simon S Du, Ken-ichi Kawarabayashi, and Stefanie Jegelka.
 696 What can neural networks reason about? *arXiv preprint arXiv:1905.13211*, 2019.
- 697
- 698 Deunsol Yoon, Hyungseok Song, Kanghoon Lee, and Woohyun Lim. CADO: Cost-aware diffusion
 699 framework for combinatorial optimization. In *ICML 2024 Workshop*, 2024.
- 700
- 701 Cong Zhang, Wen Song, Zhiguang Cao, Jie Zhang, Puay Siew Tan, and Xu Chi. Learning to
 dispatch for job shop scheduling via deep reinforcement learning. *Advances in Neural Information
 Processing Systems*, 33:1621–1632, 2020.

- 702 Lingkai Zhang, Yue Wang, and Quanquan Gu. Regularized langevin dynamics for combinatorial
703 optimization. *arXiv preprint arXiv:2502.00277*, 2025.
704
- 705 Qinsheng Zhang and Yongxin Chen. Fast sampling of diffusion models with exponential integrator.
706 *arXiv preprint arXiv:2204.13902*, 2022a.
707
- 708 Qinsheng Zhang and Yongxin Chen. Fast sampling of diffusion models with exponential integrator.
709 *arXiv preprint arXiv:2204.13902*, 2022b.
710
- 711 Zikun Zhang, Zixiang Chen, and Quanquan Gu. Convergence of score-based discrete diffusion
712 models: A discrete-time analysis. *arXiv preprint arXiv:2410.02321*, 2024.
713
- 714 Hang Zhao, Kexiong Yu, Yuhang Huang, Renjiao Yi, Chenyang Zhu, and Kai Xu. Disco: Ef-
715 ficient diffusion solver for large-scale combinatorial optimization problems. *arXiv preprint
716 arXiv:2406.19705*, 2024.
717
- 718 Kaiwen Zheng, Cheng Lu, Jianfei Chen, and Jun Zhu. Dpm-solver-v3: Improved diffusion ode
719 solver with empirical model statistics. In *Advances in Neural Information Processing Systems*,
720 volume 36, 2023.
721
- 722
- 723
- 724
- 725
- 726
- 727
- 728
- 729
- 730
- 731
- 732
- 733
- 734
- 735
- 736
- 737
- 738
- 739
- 740
- 741
- 742
- 743
- 744
- 745
- 746
- 747
- 748
- 749
- 750
- 751
- 752
- 753
- 754
- 755

APPENDICES

| | | |
|-----|---|----|
| 756 | Default Notation | 16 |
| 757 | | |
| 758 | | |
| 759 | | |
| 760 | | |
| 761 | A. Extension to Capacitated Vehicle Routing Problem (CVRP) | 17 |
| 762 | | |
| 763 | B. Proofs | 19 |
| 764 | | |
| 765 | B.1 Proof of Proposition 1 | 19 |
| 766 | B.2 Proof of E(2)-Equivariance in EDISCO | 20 |
| 767 | | |
| 768 | | |
| 769 | C. Extended Related Work | 21 |
| 770 | | |
| 771 | C.1 Foundational Neural Combinatorial Optimization | 21 |
| 772 | C.2 Alternative Architectures and Scaling Approaches | 21 |
| 773 | C.3 Discrete Diffusion Foundations and Variants | 21 |
| 774 | C.4 Theoretical Foundations and Sample Complexity | 22 |
| 775 | C.5 Hybrid and Practical Approaches | 22 |
| 776 | D. Architecture Details | 22 |
| 777 | | |
| 778 | D.1 Network Architecture Overview | 22 |
| 779 | D.2 Feature Representations and Initialization | 22 |
| 780 | D.3 Equivariant Message Passing Mechanism | 23 |
| 781 | D.4 Continuous-Time Diffusion Specifications | 23 |
| 782 | | |
| 783 | | |
| 784 | E. Additional Experiment Details | 23 |
| 785 | | |
| 786 | E.1 Performance Metrics | 23 |
| 787 | E.2 Hardware Platform | 24 |
| 788 | E.3 Results Randomness | 24 |
| 789 | E.4 Data Generation Process | 24 |
| 790 | E.5 Model Architecture Specifications | 24 |
| 791 | E.6 Training Configuration | 24 |
| 792 | F. Additional Results | 25 |
| 793 | | |
| 794 | F.1 Solver Evaluation on TSP-500 | 25 |
| 795 | F.2 TSP-10000 Results | 25 |
| 796 | F.3 TSPLIB Results | 27 |
| 797 | F.4 Noise Schedule Design for TSP Diffusion | 27 |
| 798 | F.5 Evaluation on Adaptive Mixing Parameters | 29 |
| 799 | F.6 Evaluation on Architectural Hyperparameters | 30 |
| 800 | F.7 Model Efficiency | 31 |
| 801 | F.8 Visualization of EDISCO’s E(2)-Equivariance | 32 |
| 802 | | |
| 803 | | |
| 804 | | |
| 805 | | |
| 806 | | |
| 807 | | |
| 808 | | |
| 809 | | |

810 DEFAULT NOTATION
811

812

813 NUMBERS AND ARRAYS
814

815

| | |
|-------------------------------|---|
| 816 a | A scalar (integer or real) |
| 817 \mathbf{v} | A vector |
| 818 \mathbf{M} | A matrix |
| 819 \mathcal{T} | A tensor |
| 820 \mathbf{I}_n | Identity matrix with n rows and n columns |
| 821 $\mathbf{1}$ | Vector of ones (dimensionality implied by context) |
| 822 $\mathbf{0}$ | Vector or matrix of zeros (dimensionality implied by context) |
| 823 $\text{diag}(\mathbf{v})$ | A square, diagonal matrix with diagonal entries given by \mathbf{v} |

824

825

826 GRAPH AND COMBINATORIAL STRUCTURES
827

828

| | |
|--|--|
| 829 $G = (V, E)$ | A graph with vertex set V and edge set E |
| 830 V | Set of n nodes/cities |
| 831 $\mathbf{c}_i \in \mathbb{R}^2$ | Coordinates of node/city i |
| 832 $\mathbf{X} \in \{0, 1\}^{n \times n}$ | Binary adjacency matrix representing edges |
| 833 X_{ij} | Element (i, j) of adjacency matrix (1 if edge exists, 0 otherwise) |
| 834 d_{ij} | Euclidean distance between nodes i and j |
| 835 \mathcal{C} | Constraint condition set for valid tours |

836

837

838

839 DIFFUSION PROCESS
840

841

| | |
|----------------------------------|--|
| 842 X_0 | Clean data at time $t = 0$ |
| 843 X_t | Noisy data at time $t \in [0, 1]$ |
| 844 $\beta(t)$ | Time-dependent noise schedule |
| 845 $\beta_{\min}, \beta_{\max}$ | Minimum and maximum noise rates |
| 846 $\mathbf{Q}(t)$ | Rate matrix for continuous-time Markov chain |
| 847 $\mathbf{P}(t s)$ | Transition probability matrix from time s to t |
| 848 K | Number of categorical states (2 for binary edges) |
| 849 $q(\cdot)$ | Forward diffusion distribution |
| 850 $p_\theta(\cdot)$ | Reverse diffusion distribution parameterized by θ |

851

852

853 NEURAL NETWORK COMPONENTS
854

855

| | |
|--------------------------------|---|
| 856 $\mathbf{h}_i^{(\ell)}$ | Node features for node i at layer ℓ |
| 857 $\mathbf{e}_{ij}^{(\ell)}$ | Edge features between nodes i and j at layer ℓ |
| 858 $\mathbf{x}_i^{(\ell)}$ | Coordinate embedding for node i at layer ℓ |
| 859 $\mathbf{m}_{ij}^{(\ell)}$ | Message from node j to node i at layer ℓ |
| 860 s_θ | Score network with parameters θ |
| 861 α | Step size for coordinate updates |
| 862 τ | Temperature parameter for weight scaling |
| 863 $w(t)$ | Time-dependent mixing weight function |

864

GEOMETRIC TRANSFORMATIONS

865

866

| | |
|---|--|
| 867 $E(2)$ | Euclidean group in 2D (rotations, translations, reflections) |
| 868 $SO(2)$ | Special orthogonal group (rotations) |
| 869 $g \in E(2)$ | A Euclidean transformation |
| 870 $g \cdot \mathbf{c}$ | Action of transformation g on coordinates \mathbf{c} |
| 871 \mathcal{X}/G | Quotient space under group action |
| 872 $\pi : \mathcal{X} \rightarrow \mathcal{X}/G$ | Canonical projection to quotient space |
| 873 $A \rtimes B$ | Semi-direct product of groups A and B |

874

PROBABILITY AND OPTIMIZATION

875

876

877

| | |
|-----------------------------|--|
| 878 $p(X \{\mathbf{c}_i\})$ | Conditional distribution of tours given coordinates |
| 879 $\mathbb{E}[\cdot]$ | Expectation |
| 880 $\text{Cat}(\cdot)$ | Categorical distribution |
| 881 δ_{ij} | Kronecker delta (1 if $i = j$, 0 otherwise) |
| 882 \mathcal{L} | Loss function |
| 883 NFE | Number of function evaluations |
| 884 Gap | Optimality gap: $(L_{\text{pred}} - L_{\text{opt}})/L_{\text{opt}} \times 100\%$ |

885

FUNCTIONS AND OPERATIONS

886

887

| | |
|-------------------------------|--|
| 888 $\ \cdot\ _2$ | Euclidean norm |
| 889 \oplus | Concatenation operation |
| 890 \odot | Element-wise multiplication |
| 891 \circ | Function composition, $(f \circ g)(x) = f(g(x))$ |
| 892 $\sigma(\cdot)$ | Sigmoid activation function |
| 893 $\tanh(\cdot)$ | Hyperbolic tangent activation |
| 894 $\text{MLP}(\cdot)$ | Multi-layer perceptron |
| 895 $\text{LayerNorm}(\cdot)$ | Layer normalization |
| 896 $\text{SiLU}(\cdot)$ | Sigmoid Linear Unit activation |
| 897 $\text{softmax}(\cdot)$ | Softmax function |
| 898 $\text{argmax}(\cdot)$ | Argument of the maximum |

899

A EXTENSION TO CAPACITATED VEHICLE ROUTING PROBLEM (CVRP)

900

901

902

Problem Formulation and Equivariance Preservation The Capacitated Vehicle Routing Problem (CVRP) extends TSP by introducing vehicle capacity constraints and requiring multiple routes from a central depot. While maintaining the geometric structure of TSP, CVRP presents additional challenges: (1) handling heterogeneous node types (depot vs. customers), (2) incorporating demand constraints, and (3) generating multiple feasible routes.

903

904

To preserve $E(2)$ equivariance in CVRP, we separate geometric and non-geometric features:

905

906

907

908

909

910

911

912

913

914

915

916

917

- **Equivariant features:** Customer and depot coordinates $\mathbf{c} \in \mathbb{R}^{n \times 2}$ that transform under rotations and translations
- **Invariant features:** Customer demands $d_i \in \mathbb{R}^+$ and depot indicator $\mathbf{1}_{\text{depot}}$ that remain unchanged under geometric transformations

The key insight is that while coordinates must flow through equivariant layers, demands and capacity constraints are problem-specific invariants that should not be mixed with geometric representations. Our EGNN architecture processes these separately, combining them only through invariant operations (distances and message passing).

918 **Feature Separation in EGNN Layers** Unlike TSP where node features are initialized from coordinates, CVRP initializes node embeddings exclusively from invariant features:

$$h_i^{(0)} = \text{NodeEmbed}([d_i, \mathbb{1}_{\text{depot}}(i)]) \quad (12)$$

922 where d_i is the demand of customer i and $\mathbb{1}_{\text{depot}}(i)$ indicates whether node i is the depot. This
923 ensures that geometric transformations of coordinates do not affect the initial node representations,
924 maintaining strict equivariance.

926 **Capacity-Aware Greedy Decoding** The greedy decoder for CVRP incorporates domain-specific
927 heuristics to construct feasible routes while respecting capacity constraints.

928 Edge scores are computed following the same method as in TSP: $s_{ij} = (P_{ij} + P_{ji})/d_{ij}$, where d_{ij}
929 is the Euclidean distance between nodes i and j . The symmetrization $(P_{ij} + P_{ji})$ accounts for the
930 undirected nature of the routing problem.

931 Routes are constructed iteratively while maintaining feasibility constraints. Starting with an empty
932 set of routes \mathcal{R} and unvisited customers $\mathcal{U} = \{1, \dots, n\}$, each new route begins by selecting the
933 highest-scoring feasible edge from the depot to a customer j^* whose demand d_{j^*} does not exceed
934 the vehicle capacity C . The route is then extended greedily by iteratively selecting the next customer
935 $k^* = \arg \max_{k \in \mathcal{U}} s_{jk}$ subject to the capacity constraint $\sum_{i \in \mathcal{R}_r} d_i + d_k \leq C$, where \mathcal{R}_r denotes the
936 current route being constructed. When no feasible extensions exist due to capacity limitations, the
937 vehicle returns to the depot, and a new route is initiated if unvisited customers remain.

938 To ensure solution completeness, any customers that remain unvisited after the main construction
939 phase are assigned to individual routes. This post-processing step guarantees that all customers
940 are served, though it may result in suboptimal routing for instances with tight capacity constraints.
941 The overall approach balances solution quality with computational efficiency while maintaining the
942 geometric structure learned by the equivariant network.

943 We evaluate EDISCO on standard CVRP benchmarks with 20, 50, and 100 customers, following
944 the evaluation protocol from Kool et al. (2019). All instances use a vehicle capacity of 50 units with
945 customer demands uniformly sampled from $\{1, \dots, 9\}$.

947 Table 4: Results on CVRP-20, CVRP-50, and CVRP-100. RL: Reinforcement Learning, SL: Su-
948 pervised Learning, G: Greedy, S: Sampling. LKH-3* represents the baseline for computing the gap.
949 Gurobi results are from the user’s provided table. AM results are from Kool et al. (2019). POMO
950 results are from Kwon et al. (2020). Sym-NCO results are from Kim et al. (2022) (CVRP-100 only).

| 952 Algorithm | 953 Type | 954 CVRP-20 | | 955 CVRP-50 | | 956 CVRP-100 | |
|---|-----------------|--------------------|--------------|--------------------|--------------|---------------------|--------------|
| | | 957 Cost↓ | 958 Gap↓ | 959 Cost↓ | 960 Gap↓ | 961 Cost↓ | 962 Gap↓ |
| Gurobi (Gurobi Optimization, LLC, 2020) | Exact | 6.10 | 0.00% | — | — | — | — |
| LKH-3* (Helsgaun, 2017) | Heuristic | 6.14 | 0.58% | 10.38 | 0.00% | 15.65 | 0.00% |
| <i>Greedy Decoding</i> | | | | | | | |
| AM (Kool et al., 2019) | RL+G | 6.40 | 4.97% | 10.98 | 5.86% | 16.80 | 7.34% |
| POMO (Kwon et al., 2020) | RL+G | 6.35 | 3.72% | 10.74 | 3.52% | 16.13 | 3.09% |
| Sym-NCO (Kim et al., 2022) | RL+G | — | — | — | — | 16.10 | 2.88% |
| EDISCO (ours) | SL+G | 6.21 | 1.41% | 10.63 | 2.46% | 16.15 | 3.17% |
| <i>Sampling/Multiple Trajectories</i> | | | | | | | |
| AM (Kool et al., 2019) | RL+S (1280) | 6.25 | 2.49% | 10.62 | 2.40% | 16.23 | 3.72% |
| POMO (no aug) (Kwon et al., 2020) | RL+G | 6.17 | 0.82% | 10.49 | 1.14% | 15.83 | 1.13% |
| POMO (x8 aug) (Kwon et al., 2020) | RL+G | 6.14 | 0.21% | 10.42 | 0.45% | 15.73 | 0.51% |
| Sym-NCO (Kim et al., 2022) | RL+S (100) | — | — | — | — | 15.87 | 1.40% |
| EDISCO (ours) | SL+S (1280) | 6.15 | 0.33% | 10.41 | 0.35% | 15.71 | 0.38% |

966 **Analysis of Results** The results are shown in Table 4. EDISCO demonstrates strong performance
967 on CVRP, achieving competitive results across all instance sizes. With greedy decoding, EDISCO
968 substantially outperforms AM (1.41% vs 4.97% on CVRP-20), though it slightly trails Sym-NCO
969 on CVRP-100 (3.17% vs 2.88%). With sampling, EDISCO achieves the best performance on larger
970 instances (0.35% on CVRP-50, 0.38% on CVRP-100) but falls behind POMO with augmentation
971 on CVRP-20 (0.33% vs 0.21%).

These results highlight fundamental differences between diffusion models and autoregressive approaches. Autoregressive models like POMO generate solutions sequentially, enabling them to leverage explicit augmentation strategies that evaluate problems from multiple geometric perspectives. POMO’s 8x augmentation effectively multiplies the search space by considering rotated and reflected problem instances. In contrast, EDISCO’s diffusion process operates on entire adjacency matrices simultaneously, learning to denoise complete solutions rather than constructing them step-by-step. While this parallel generation captures global solution patterns more effectively, it cannot directly benefit from the same augmentation multipliers that boost autoregressive performance on small instances.

B PROOFS

B.1 PROOF OF PROPOSITION 1

Proof. We establish each claim systematically.

Part (i): Quotient manifold structure. The Euclidean group $E(2) = \mathbb{R}^2 \rtimes SO(2)$ is a 3-dimensional Lie group, where \mathbb{R}^2 corresponds to translations and $SO(2)$ to rotations. The action of $g = (t, R) \in E(2)$ on a configuration $\mathbf{x} = (x_1, y_1, \dots, x_n, y_n) \in X$ is given by:

$$g \cdot \mathbf{x} = (Rx'_1 + t, Ry'_1 + t, \dots, Rx'_n + t, Ry'_n + t)$$

where $x'_i = (x_i, y_i)^T$ denotes the i -th city as a column vector, and $R \in SO(2)$ acts by matrix multiplication.

To show the action is free, suppose $g \cdot \mathbf{x} = \mathbf{x}$ for some $g = (t, R) \in G$ and $\mathbf{x} \in X$. This means:

$$R \begin{pmatrix} x_i \\ y_i \end{pmatrix} + t = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \quad \forall i \in \{1, \dots, n\}$$

For $i = 1$, we have $(R - I) \begin{pmatrix} x_1 \\ y_1 \end{pmatrix} = -t$. For $i = 2$, we have $(R - I) \begin{pmatrix} x_2 \\ y_2 \end{pmatrix} = -t$. Subtracting these equations:

$$(R - I) \begin{pmatrix} x_1 - x_2 \\ y_1 - y_2 \end{pmatrix} = 0$$

If the cities are not all identical (which we assume for any meaningful TSP instance), there exists at least one pair (i, j) such that $(x_i - x_j, y_i - y_j) \neq (0, 0)$. For $R \neq I$, the matrix $(R - I)$ has rank 2 (its eigenvalues are $e^{i\theta} - 1$ and $e^{-i\theta} - 1$ for rotation angle $\theta \neq 0$), so it has trivial kernel. Thus $(R - I)v = 0$ for non-zero v implies $R = I$.

With $R = I$, the condition becomes $t = 0$, so $g = e$ is the identity. Therefore, the action is free.

By the quotient manifold theorem (Lee, 2012), when a Lie group G acts freely and properly on a manifold M , the quotient M/G inherits a unique smooth manifold structure such that $\pi : M \rightarrow M/G$ is a smooth submersion. The dimension formula gives:

$$\dim(X/G) = \dim(X) - \dim(G) = 2n - 3$$

Part (ii): Factorization of equivariant functions. Let $F : X \rightarrow Y$ be a G -equivariant function, meaning $F(g \cdot \mathbf{x}) = \rho(g) \cdot F(\mathbf{x})$ for all $g \in G$ and $\mathbf{x} \in X$, where $\rho : G \rightarrow \text{Aut}(Y)$ is a representation of G on Y .

For TSP edge prediction, we typically have $Y = \{0, 1\}^{n \times n}$ (adjacency matrices) and ρ is the trivial representation (since the optimal tour is invariant under Euclidean transformations of the cities). Thus $F(g \cdot \mathbf{x}) = F(\mathbf{x})$ for all $g \in G$.

Define $\tilde{F} : X/G \rightarrow Y$ by $\tilde{F}([\mathbf{x}]) = F(\mathbf{x})$, where $[\mathbf{x}] = \{g \cdot \mathbf{x} : g \in G\}$ denotes the orbit of \mathbf{x} . This is well-defined precisely because of equivariance: if $[\mathbf{x}] = [\mathbf{x}']$, then $\mathbf{x}' = g \cdot \mathbf{x}$ for some $g \in G$, so:

$$\tilde{F}([\mathbf{x}']) = F(\mathbf{x}') = F(g \cdot \mathbf{x}) = F(\mathbf{x}) = \tilde{F}([\mathbf{x}])$$

1026 The factorization $F = \tilde{F} \circ \pi$ follows immediately from the definition:
 1027
 1028

$$F(\mathbf{x}) = \tilde{F}([\mathbf{x}]) = \tilde{F}(\pi(\mathbf{x})) = (\tilde{F} \circ \pi)(\mathbf{x})$$

1029 Uniqueness of \tilde{F} follows from surjectivity of π : if $F = \tilde{F}_1 \circ \pi = \tilde{F}_2 \circ \pi$, then for any $[\mathbf{x}] \in X/G$,
 1030 choosing any representative $\mathbf{x} \in [\mathbf{x}]$:

$$\tilde{F}_1([\mathbf{x}]) = \tilde{F}_1(\pi(\mathbf{x})) = F(\mathbf{x}) = \tilde{F}_2(\pi(\mathbf{x})) = \tilde{F}_2([\mathbf{x}])$$

1032
 1033
 1034 **Part (iii): Learning complexity reduction.** The factorization $F = \tilde{F} \circ \pi$ establishes a bijection
 1035 between:

$$\begin{aligned} \{G\text{-equivariant functions } X \rightarrow Y\} &\longleftrightarrow \{\text{functions } X/G \rightarrow Y\} \\ F &\longmapsto \tilde{F} \\ \tilde{F} \circ \pi F & \end{aligned}$$

1036 Therefore, learning any G -equivariant function F is equivalent to learning the corresponding func-
 1037 tion \tilde{F} on the quotient manifold. Since X/G has dimension $2n - 3$ while X has dimension $2n$, the
 1038 domain of \tilde{F} has three fewer degrees of freedom.

1039 In terms of function approximation, this means:

- 1040 • A basis of functions on X/G requires parametrization by $2n - 3$ variables
- 1041 • Local charts for X/G have dimension $2n - 3$
- 1042 • The metric entropy and covering numbers scale with the intrinsic dimension $2n - 3$

1043 This completes the proof that E(2)-equivariant learning reduces to learning on a lower-dimensional
 1044 manifold, providing the theoretical foundation for improved sample efficiency. \square
 1045

1046 B.2 PROOF OF E(2)-EQUIVARIANCE IN EDISCO

1047 *Proof.* We prove that E(2)-equivariance is preserved throughout the entire EDISCO pipeline, from
 1048 input processing through diffusion to tour construction.

1049 **Step 1: EGNN Architecture Preserves Equivariance**

1050 For any Euclidean transformation $g \in E(n)$, we show each layer maintains equivariance:

1051 (i) *Distance invariance:* For transformed coordinates $g \cdot c_i$:

$$d_{ij}^{(g)} = \|g \cdot c_i - g \cdot c_j\|_2 = \|g(c_i - c_j)\|_2 = \|c_i - c_j\|_2 = d_{ij}$$

1052 (ii) *Message invariance:* From Equation 8, messages depend only on: - Node features h_i, h_j (ini-
 1053 tialized as invariant city indices) - Edge features e_{ij} (initialized from noisy adjacency matrix) -
 1054 Distances d_{ij} (proven invariant above)

1055 Therefore: $m_{ij}^{(g)} = m_{ij}$

1056 (iii) *Coordinate equivariance:* The update rule (Equation 9):

$$\Delta(g \cdot x_i) = \alpha \sum_{j \neq i} w_{ij} \cdot \frac{g \cdot x_j - g \cdot x_i}{\|g \cdot x_j - g \cdot x_i\|_2} \quad (13)$$

$$= \alpha \sum_{j \neq i} w_{ij} \cdot \frac{g(x_j - x_i)}{\|x_j - x_i\|_2} \quad (14)$$

$$= g \cdot \left(\alpha \sum_{j \neq i} w_{ij} \cdot \frac{x_j - x_i}{\|x_j - x_i\|_2} \right) \quad (15)$$

$$= g \cdot \Delta x_i \quad (16)$$

1080
 1081 (iv) *Feature invariance:* Edge and node feature updates depend only on invariant quantities, thus
 1082 $e_{ij}^{(g)} = e_{ij}$ and $h_i^{(g)} = h_i$.

1083 **Step 2: Diffusion Process Maintains Equivariance**

1084 The categorical diffusion operates on edge variables $X_{ij} \in \{0, 1\}$, which represent whether edge
 1085 (i, j) is in the tour. These are inherently invariant to coordinate transformations.

1086 *Forward process:* The corruption adds noise to edge selections independent of coordinates:

1088
$$q(X_t | X_0) = \prod_{i,j} \text{Cat}(X_{t,ij} | p = P_{ij}(t|0))$$

 1089

1090 *Reverse process:* Since the score network s_θ outputs edge probabilities that are invariant (proven in
 1091 Step 1), the reverse process maintains this invariance:

1093
$$p_\theta(X_{t-\Delta t} | X_t, g(\{c_i\})) = p_\theta(X_{t-\Delta t} | X_t, \{c_i\})$$

1094 **Step 3: Tour Construction Preserves Optimality**

1095 The greedy decoding computes scores:

1097
$$s_{ij}^{(g)} = \frac{P_{ij} + P_{ji}}{d_{ij}^{(g)}} = \frac{P_{ij} + P_{ji}}{d_{ij}} = s_{ij}$$

 1098
 1099

1100 Since edge scores are identical under transformation, the greedy algorithm produces tours with
 1101 identical edge selections (up to vertex relabeling). \square

1103 **C EXTENDED RELATED WORK**

1105 **C.1 FOUNDATIONAL NEURAL COMBINATORIAL OPTIMIZATION**

1107 The application of neural networks to combinatorial optimization began with Pointer Networks
 1108 (Vinyals et al., 2015), which introduced attention mechanisms to construct variable-
 1109 length permutations. While this required supervised training with optimal solutions, subsequent
 1110 work (Bello et al., 2017) demonstrated that reinforcement learning could discover effective heuris-
 1111 tics without labeled data, eliminating a major practical limitation. The evolution continued with
 1112 the attention model (Kool et al., 2019), which improved upon Pointer Networks through multi-head
 1113 attention and achieved strong performance without problem-specific design. POMO (Kwon et al.,
 1114 2020) further advanced autoregressive methods by exploring multiple rollouts from different starting
 1115 points. These foundational works established that neural networks could learn meaningful represen-
 1116 tations of combinatorial structure, though they struggled with generalization to larger instances (Fu
 1117 et al., 2021).

1118 **C.2 ALTERNATIVE ARCHITECTURES AND SCALING APPROACHES**

1119 Beyond diffusion-based methods, several innovative architectures address the challenge of scal-
 1120 ing to large TSP instances. LEHD (Light Encoder Heavy Decoder) (Fu et al., 2024) achieves
 1121 remarkable scalability to instances with up to 10,000 cities by separating encoding and decoding
 1122 complexity—training on small instances but generalizing through architectural design rather than
 1123 data. Bisimulation quotienting (BQ-NCO) (Drakulic et al., 2023) takes a fundamentally different
 1124 approach by reformulating the MDP to group behaviorally similar states, achieving strong zero-shot
 1125 generalization. Hierarchical approaches like GLOP (Corsini et al., 2024) combine global partition
 1126 with local construction for real-time routing, while the hierarchical neural constructive solver (Goh
 1127 et al., 2024) builds solutions through multiple resolution levels. These methods demonstrate that
 1128 architectural innovations can sometimes overcome the data requirements that limit standard ap-
 1129 proaches.

1130 **C.3 DISCRETE DIFFUSION FOUNDATIONS AND VARIANTS**

1131 The theoretical foundations for discrete diffusion (Austin et al., 2021) established how to apply dif-
 1132 fusion processes to categorical data through transition matrices, providing the basis for subsequent

TSP solvers. Recent advances include variational flow matching (Akhound-Sadegh et al., 2024) and discrete flow matching (Campbell et al., 2024a), which provide alternative formulations with improved training dynamics. The comprehensive treatment of continuous diffusion for categorical data (Dieleman et al., 2022) addressed many technical details necessary for practical implementation. DeFoG (Campbell et al., 2024b) demonstrates state-of-the-art performance on graph generation through discrete flow matching, suggesting potential applications to optimization. The connection to optimal transport (Lipman et al., 2022) offers theoretical insights that could lead to algorithmic improvements, while regularized Langevin dynamics (Zhang et al., 2025) shows how continuous-time formulations avoid local optima more effectively than discrete-time approaches.

C.4 THEORETICAL FOUNDATIONS AND SAMPLE COMPLEXITY

Understanding why certain neural architectures succeed at combinatorial optimization remains an active area of research. The analysis of graph neural network expressiveness (Xu et al., 2019) establishes fundamental representation limits, while work on algorithmic alignment (Xu et al., 2021) shows that architectures matching problem structure generalize better. Recent theoretical advances prove that equivariant models achieve exponentially better sample complexity than non-equivariant ones (Brehmer et al., 2024), providing a rigorous justification for geometric inductive biases. The analysis of learning TSP and generalization (Joshi et al., 2022) demonstrates fundamental limitations of supervised approaches and suggests that architectural innovations are necessary for progress. Convergence analysis for discrete diffusion models (Zhang et al., 2024) provides rates that inform practical algorithm design, while the study of instance hardness (Smith-Miles et al., 2010) reveals what makes problems difficult for neural solvers.

C.5 HYBRID AND PRACTICAL APPROACHES

Combining neural networks with classical optimization algorithms leverages complementary strengths. Learning to perform local rewriting (Chen & Tian, 2019) trains networks to improve existing solutions through targeted modifications, while integration with branch-and-bound (Gasse et al., 2019) accelerates exact algorithms through learned branching strategies. Neural diving (Nair et al., 2020) combines neural networks with MIP solvers for fast feasible solution finding. These hybrid methods often outperform purely neural or classical approaches, suggesting that practical deployment may require combining paradigms. Recent work on unsupervised learning (Wang & Li, 2023) and self-improvement (Hudson et al., 2024) reduces dependence on high-quality training data, addressing a major practical limitation. Applications beyond TSP demonstrate broader impact, including vehicle routing with complex constraints (Nazari et al., 2018), scheduling (Zhang et al., 2020), and circuit design (Mirhoseini et al., 2021).

D ARCHITECTURE DETAILS

D.1 NETWORK ARCHITECTURE OVERVIEW

The EDISCO model employs a 12-layer E(2)-equivariant graph neural network that processes city coordinates and noisy adjacency matrices while maintaining geometric equivariance. The architecture consists of three main components: an embedding module, stacked equivariant layers, and a prediction head.

D.2 FEATURE REPRESENTATIONS AND INITIALIZATION

The model maintains three distinct feature types throughout the network:

Spatial Features. City coordinates $\mathbf{c} \in \mathbb{R}^{n \times 2}$ are transformed into 64-dimensional node embeddings via a linear projection. Additionally, coordinate representations $\mathbf{x} \in \mathbb{R}^{n \times 2}$ are maintained separately and evolve through equivariant updates during message passing.

Relational Features. Edge features $\mathbf{e} \in \mathbb{R}^{n \times n \times 64}$ encode pairwise relationships and tour decisions. These are initialized from the noisy adjacency matrix X_t through a single linear transformation.

1188
 1189 **Temporal Encoding.** The continuous diffusion time $t \in [0, 1]$ is encoded using sinusoidal basis
 1190 functions with frequencies spanning multiple octaves, producing a 128-dimensional representation
 1191 that modulates the network’s behavior at different noise levels.
 1192

1193 D.3 EQUIVARIANT MESSAGE PASSING MECHANISM

1194 Each EGNN layer performs the following operations while preserving E(n) symmetry:
 1195

1196 **Message Formation.** Pairwise messages aggregate local and geometric information:

$$1197 \quad \mathbf{m}_{ij} = f_{\text{msg}}(\mathbf{h}_i \oplus \mathbf{h}_j \oplus \mathbf{e}_{ij} \oplus d_{ij})$$

1199 where $d_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2$ provides rotation-invariant distance information and f_{msg} is a 3-layer MLP
 1200 with SiLU activations and layer normalization.
 1201

1202 **Geometric Updates.** Coordinate evolution respects equivariance constraints through normalized
 1203 directional updates:
 1204

$$1205 \quad \mathbf{x}_i \leftarrow \mathbf{x}_i + 0.1 \sum_j \text{Gate}(\mathbf{m}_{ij}) \cdot \frac{\mathbf{x}_j - \mathbf{x}_i}{\|\mathbf{x}_j - \mathbf{x}_i\|_2 + 10^{-8}}$$

1207 The gating function employs a temperature-scaled tanh with $\tau = 10$ to prevent gradient saturation.
 1208

1209 **Feature Evolution.** Node and edge features incorporate aggregated messages through residual
 1210 connections:
 1211

$$1212 \quad \mathbf{h}_i \leftarrow \text{LN}(\mathbf{h}_i + f_{\text{node}}(\mathbf{h}_i, \sum_j \mathbf{m}_{ij})) \quad (17)$$

$$1214 \quad \mathbf{e}_{ij} \leftarrow \text{LN}(\mathbf{e}_{ij} + f_{\text{edge}}(\mathbf{e}_{ij}, \mathbf{m}_{ij}) + f_{\text{time}}(\mathbf{t})) \quad (18)$$

1215 where LN denotes layer normalization and f_{node} , f_{edge} , f_{time} are learned transformations.
 1216

1217 D.4 CONTINUOUS-TIME DIFFUSION SPECIFICATIONS

1219 **Forward Process.** The categorical diffusion operates on binary edge variables through a
 1220 continuous-time Markov chain with rate matrix:
 1221

$$1222 \quad Q(t) = \beta(t) \begin{bmatrix} -0.5 & 0.5 \\ 0.5 & -0.5 \end{bmatrix}$$

1224 where $\beta(t)$ increases linearly from 0.1 to 1.5 over the unit interval.
 1225

1226 **Transition Dynamics.** The forward transition probability admits a closed-form solution:
 1227

$$1228 \quad P(X_t = j | X_0 = i) = \frac{1}{2} + \left(\delta_{ij} - \frac{1}{2} \right) \exp \left(-2 \int_0^t \beta(u) du \right)$$

1230 **Reverse Sampling.** The model employs an adaptive mixing strategy that interpolates between
 1231 diffusion dynamics and direct prediction:
 1232

- 1233 • For $t > 0.1$: Stochastic transitions weighted by $w(t) = t$.
 1234
- 1235 • For $t \leq 0.1$: Deterministic argmax selection.
 1236
- 1237 • Default sampling uses 50 steps with optional adaptive scheduling.

1238 E ADDITIONAL EXPERIMENT DETAILS

1240 E.1 PERFORMANCE METRICS

1241 We evaluate models using three criteria:

- **Tour Length:** Average Euclidean length of generated tours across test instances
- **Optimality Gap:** Relative deviation from optimal/best-known solutions, computed as $(L_{\text{model}} - L_{\text{optimal}})/L_{\text{optimal}} \times 100\%$
- **Inference Duration:** Wall-clock time for generating solutions on the test set, measured in seconds (s) or minutes (m)

E.2 HARDWARE PLATFORM

All experiments were conducted on a single NVIDIA RTX A6000 GPU paired with dual Intel Xeon Gold 5218R CPUs. Both training and inference use the same hardware configuration.

E.3 RESULTS RANDOMNESS

Due to the stochastic nature of diffusion models, all results reported are the averaged results over five runs with different random seeds.

E.4 DATA GENERATION PROCESS

Instance Creation We follow the exact data generation protocol from DIFUSCO (Sun & Yang, 2023) for fair comparison. All cities are sampled uniformly from the unit square $[0, 1]^2$ following standard practice in the TSP literature. For smaller instances, TSP-50 and TSP-100 problems are solved to optimality using the Concorde exact solver (Applegate et al., 2006) to obtain ground truth tours. For larger scales, TSP-500 and TSP-1000 instances are labeled using the LKH-3 heuristic solver (Helsgaun, 2017) with 500 trials to ensure near-optimal solution quality. Our evaluation employs the standard test sets from Kool et al. (Kool et al., 2019) for TSP-50/100 containing 1,280 instances each, and from Fu et al. (Fu et al., 2021) for TSP-500/1000 containing 128 instances each.

Graph Sparsification For problems exceeding 100 cities, computational efficiency necessitates graph sparsification strategies. We implement k-nearest neighbor sparsification where each city connects only to its k closest neighbors based on Euclidean distance, setting k=50 for TSP-500 and k=100 for TSP-1000. This distance-based edge pruning dramatically reduces the computational complexity from $O(n^2)$ to $O(nk)$ while preserving the most relevant edges for tour construction. Correspondingly, dense matrix operations are replaced with their sparse equivalents throughout the network architecture to maintain computational efficiency at scale.

E.5 MODEL ARCHITECTURE SPECIFICATIONS

For all experiments, the network contains approximately 5.5M trainable parameters distributed across:

- 12 EGNN layers with shared architecture
- Node dimension: 64
- Edge dimension: 64
- Hidden dimension: 256
- Timestep embedding dimension: 128

This setting ensures that EDISCO has a similar number of trainable parameters to the SOTA diffusion TSP solvers (5.3M) (Sun & Yang, 2023; Li et al., 2023; Yoon et al., 2024; Zhao et al., 2024), allowing for a fair comparison.

E.6 TRAINING CONFIGURATION

We train EDISCO using the AdamW optimizer with a learning rate of 2×10^{-4} and weight decay of 10^{-5} . The learning rate follows a cosine annealing schedule over the training epochs to ensure smooth convergence. For training stability, we apply gradient clipping at unit norm to prevent exploding gradients during the reverse diffusion process. The loss function employs a simplified

ELBO formulation with time-dependent weighting $(1 - \sqrt{t})$, which emphasizes reconstruction accuracy near $t = 0$ while maintaining stable gradients throughout the diffusion trajectory.

- **TSP-50:** 500,000 training instances, batch size 64, 50 epochs.
- **TSP-100:** 500,000 training instances, batch size 32, 50 epochs.
- **TSP-500:** 60,000 instances, batch size 16, 50 epochs with curriculum learning initialized from TSP-100 checkpoint.
- **TSP-1000:** 30,000 instances, batch size 8, 50 epochs with curriculum learning initialized from TSP-100 checkpoint.
- **TSP-10000:** 3,000 instances, batch size 4, 50 epochs with curriculum learning initialized from TSP-500 checkpoint.

F ADDITIONAL RESULTS

F.1 SOLVER EVALUATION ON TSP-500

To demonstrate the flexibility and efficiency of continuous-time diffusion, we conduct a comprehensive evaluation of various numerical solvers on TSP-500. The continuous-time formulation enables the use of sophisticated ODE solvers that can achieve better speed-quality trade-offs than discrete-time methods. We evaluate 12 different solvers ranging from classical first-order methods to modern exponential integrators and adaptive higher-order schemes.

Table 5 presents the results across different solver families. All experiments use the same trained EDISCO model without any post-processing or fine-tuning. Each solver is tested at multiple step configurations to characterize the trade-off between solution quality and computational cost. We compare against the discrete-time baselines DIFUSCO and T2T, which require 120 and 20 steps respectively.

The results reveal several key findings. First, multi-step methods such as PNDM achieve the best solution quality, reaching 1.95% optimality gap with 50 steps (51 NFE) in 2.19 minutes. This represents a 2.6x speedup over DIFUSCO (5.70m) while achieving substantially better solution quality (1.95% vs 9.41% gap). Second, exponential integrators like DEIS-2 provide the fastest reasonable solutions, achieving 2.78% gap in only 0.23 minutes with 5 steps. This 25x speedup over DIFUSCO demonstrates the practical advantages of continuous-time formulation for real-time applications.

Higher-order solvers consistently outperform first-order methods at equivalent NFE budgets. For instance, Heun’s method (RK2) achieves 1.99% gap with 20 NFE in 0.83 minutes, while the first-order Euler method reaches only 3.14% gap with 11 NFE in 0.45 minutes. The classical RK4 method achieves near-optimal performance (1.97% gap) with just 5 integration steps in 0.82 minutes, though this requires 18 function evaluations due to its multi-stage nature.

Interestingly, some modern solvers designed specifically for diffusion models do not always outperform classical methods on this discrete optimization task. EDM-Heun, despite its success in image generation, produces 15.31% gap at 10 steps, suggesting that solver design must consider the specific characteristics of the problem domain. Similarly, DDIM shows poor performance (14.48% gap) compared to other first-order methods, likely due to its parameterization being optimized for continuous rather than discrete state spaces.

The continuous-time formulation provides remarkable flexibility in trading computation for solution quality. Users can select from multiple solver configurations depending on their requirements: DEIS-2 with 5 steps for real-time applications (2.78% gap, 0.23m), DPM-Solver-2 with 25 steps for balanced performance (2.03% gap, 2.33m), or PNDM with 50 steps for best quality (1.95% gap, 2.19m). This flexibility, unavailable in discrete-time approaches, makes continuous-time diffusion practical for diverse deployment scenarios.

F.2 TSP-10000 RESULTS

Table 6 presents results on the challenging TSP-10000 benchmark, demonstrating EDISCO’s scalability to very large problem instances. With greedy decoding, EDISCO achieves a 1.98% optimality

1350

1351 Table 5: Comprehensive solver evaluation on TSP-500. G: Greedy Decoding. Best gap: PNDM
1352 with 50 steps (1.95%). Fastest <3% gap: DEIS-2 with 5 steps (2.78%, 0.23m).

| Method | Type | Steps | | Performance | | Time (minutes) ↓ |
|---|-------|-------|-----|-------------|--------------|---------------------|
| | | Steps | NFE | Length ↓ | Gap ↓ | |
| Concorde* (Applegate et al., 2006) | Exact | - | - | 16.55 | 0.00% | - |
| <i>Discrete-Time Baselines</i> | | | | | | |
| DIFUSCO (Sun & Yang, 2023) | SL+G | 120 | 120 | 18.11 | 9.41% | 5.70 |
| T2T (Li et al., 2023) | SL+G | 20 | ~60 | 17.39 | 5.09% | 4.90 |
| <i>First-Order Solvers</i> | | | | | | |
| EDISCO (Euler) (Särkkä & Solin, 2019) | SL+G | 10 | 11 | 17.07 | 3.14% | 0.45 |
| EDISCO (Euler) | SL+G | 25 | 26 | 17.10 | 3.32% | 1.09 |
| EDISCO (Euler) | SL+G | 50 | 51 | 17.08 | 3.18% | 2.15 |
| EDISCO (Euler) | SL+G | 100 | 101 | 17.02 | 2.81% | 4.25 |
| EDISCO (DDIM, $\eta=0$) (Song et al., 2021a) | SL+G | 10 | 11 | 18.97 | 14.48% | 0.45 |
| EDISCO (DDIM, $\eta=0$) | SL+G | 50 | 51 | 19.35 | 17.61% | 2.13 |
| EDISCO (DDIM, $\eta=0.5$) | SL+G | 10 | 11 | 18.03 | 9.52% | 0.44 |
| <i>Multi-Step Methods</i> | | | | | | |
| EDISCO (PNDM) (Liu et al., 2022) | SL+G | 5 | 6 | 17.41 | 5.29% | 0.23 |
| EDISCO (PNDM) | SL+G | 10 | 11 | 17.05 | 3.02% | 0.43 |
| EDISCO (PNDM) | SL+G | 25 | 26 | 17.16 | 3.68% | 1.10 |
| EDISCO (PNDM) | SL+G | 50 | 51 | 16.87 | 1.95% | 2.19 |
| EDISCO (PNDM) | SL+G | 100 | 101 | 16.89 | 2.31% | 4.35 |
| <i>Exponential Integrators</i> | | | | | | |
| EDISCO (DEIS-2) (Zhang & Chen, 2022a) | SL+G | 5 | 6 | 17.01 | 2.78% | 0.23 |
| EDISCO (DEIS-2) | SL+G | 10 | 11 | 17.73 | 7.12% | 0.42 |
| EDISCO (DEIS-2) | SL+G | 25 | 26 | 18.58 | 12.26% | 1.09 |
| EDISCO (DEIS-3) | SL+G | 5 | 6 | 17.29 | 4.48% | 0.23 |
| EDISCO (DEIS-3) | SL+G | 10 | 11 | 18.31 | 10.63% | 0.45 |
| <i>Higher-Order Solvers</i> | | | | | | |
| EDISCO (Heun/RK2) (Butcher, 2016) | SL+G | 5 | 10 | 16.89 | 2.34% | 0.40 |
| EDISCO (Heun/RK2) | SL+G | 10 | 20 | 16.88 | 1.99% | 0.83 |
| EDISCO (Heun/RK2) | SL+G | 25 | 50 | 16.90 | 2.17% | 2.09 |
| EDISCO (DPM-Solver-2) (Lu et al., 2022a) | SL+G | 5 | 10 | 17.09 | 3.31% | 0.44 |
| EDISCO (DPM-Solver-2) | SL+G | 10 | 20 | 16.89 | 2.32% | 0.91 |
| EDISCO (DPM-Solver-2) | SL+G | 25 | 50 | 16.88 | 2.03% | 2.33 |
| EDISCO (DPM-Solver++) (Lu et al., 2022b) | SL+G | 5 | 6 | 17.91 | 8.26% | 0.22 |
| EDISCO (DPM-Solver++) | SL+G | 10 | 11 | 18.88 | 13.42% | 0.45 |
| EDISCO (DPM-Solver++) | SL+G | 25 | 26 | 17.01 | 2.71% | 1.11 |
| EDISCO (DPM-Solver-3) (Zheng et al., 2023) | SL+G | 5 | 14 | 16.96 | 2.48% | 0.64 |
| EDISCO (DPM-Solver-3) | SL+G | 10 | 29 | 16.95 | 2.41% | 1.35 |
| EDISCO (RK4) (Butcher, 2016) | SL+G | 5 | 18 | 16.88 | 1.97% | 0.82 |
| EDISCO (RK4) | SL+G | 10 | 38 | 16.89 | 2.13% | 1.78 |
| EDISCO (EDM-Heun) (Karras et al., 2022) | SL+G | 10 | 19 | 18.75 | 15.31% | 0.79 |
| EDISCO (EDM-Heun) | SL+G | 25 | 46 | 18.16 | 9.72% | 1.97 |

1390

1391

1392

1393

1394

1395

1396

1397 gap in 12.18 minutes, significantly outperforming DIFUSCO (8.95%, 28.51m) and surpassing both
1398 T2T (2.92%, 1.52h) and DISCO (2.90%, 1.52h) while being 7.5× faster than T2T.1399 When enhanced with 2-opt post-processing, EDISCO achieves near-optimal solutions with only
1400 0.51% gap in 12.72 minutes. Under the sampling-based decoding, EDISCO achieves a 1.79% gap
1401 compared to DIFUSCO’s 33.09% and matches the performance of T2T and DISCO (2.84%) while
1402 being 3.8× faster than T2T. With sampling plus 2-opt, EDISCO reaches an exceptional 0.20% gap in
1403 39.28 minutes, compared to DIFUSCO’s 4.03% in 6.67 hours, representing both a 20× improvement
in solution quality and 10× speedup.

Table 6: Results on TSP-10000. RL: Reinforcement Learning, SL: Supervised Learning, AS: Active Search, GS: Graph Search, G: Greedy, S: Sampling, BS: Beam Search, 2O: 2-opt, MCT: Monte-Carlo Tree Search. LKH-3 (default)* represents the baseline for computing the gap. Results for DIFUSCO are from Sun & Yang (2023). Results for DISCO, AM, and GLOP are from Zhao et al. (2024). Results for T2T are from Li et al. (2023). Results for DIMES are from Qiu et al. (2022).

| Algorithm | Type | Length↓ | Gap↓ | Time |
|--|------------|--------------|--------------|--------|
| LKH-3 (default)* (Helsgaun, 2017) | Heuristics | 71.77 | — | 8.8h |
| LKH-3 (less trials) (Helsgaun, 2017) | Heuristics | 71.79 | 0.03% | 51.27m |
| 2-opt (Lin & Kernighan, 1973) | Heuristics | 91.16 | 27.02% | 28.49m |
| Farthest Insertion | Heuristics | 80.59 | 12.36% | 13.25m |
| AM (Kool et al., 2019) | RL+G | 141.51 | 97.17% | 7.68m |
| GLOP (Corsini et al., 2023) | RL+G | 75.29 | 4.90% | 1.90m |
| DIMES (Qiu et al., 2022) | RL+AS+G | 80.45 | 12.09% | 3.07h |
| DIFUSCO (Sun & Yang, 2023) | SL+G | 78.35 | 8.95% | 28.51m |
| T2T (Li et al., 2023) | SL+G | 73.87 | 2.92% | 1.52h |
| DISCO (Zhao et al., 2024) | SL+G | 73.85 | 2.90% | 1.52h |
| EDISCO with 50-step PNDM (ours) | SL+G | 73.19 | 1.98% | 12.18m |
| DIFUSCO (Sun & Yang, 2023) | SL+G+2O | 73.99 | 3.10% | 35.38m |
| EDISCO with 50-step PNDM (ours) | SL+G+2O | 72.87 | 1.53% | 12.72m |
| AM (Kool et al., 2019) | RL+BS | 129.40 | 80.28% | 1.81h |
| GLOP (Corsini et al., 2023) | RL+S | 75.27 | 4.88% | 5.96m |
| DIFUSCO (Sun & Yang, 2023) | SL+S | 95.52 | 33.09% | 6.59h |
| T2T (Li et al., 2023) | SL+S | 73.81 | 2.84% | 2.47h |
| DISCO (Zhao et al., 2024) | SL+S | 73.81 | 2.84% | 48.77m |
| EDISCO with 50-step PNDM (ours) | SL+S | 72.77 | 1.39% | 38.92m |
| DIFUSCO (Sun & Yang, 2023) | SL+S+2O | 74.66 | 4.03% | 6.67h |
| DISCO (Zhao et al., 2024) | SL+GS+MCTS | 73.69 | 2.68% | 2.1h |
| EDISCO with 50-step PNDM (ours) | SL+S+2O | 72.63 | 1.20% | 39.28m |

F.3 TSPLIB RESULTS

We evaluate EDISCO on real-world TSP instances from the TSPLIB benchmark (Reinelt, 1991). Following prior work (Fu et al., 2021; Li et al., 2023), we train EDISCO on randomly generated 100-node TSP instances and evaluate them on TSPLIB instances ranging from 50 to 200 nodes.

Table 7 presents the optimality gaps (percentage above the known optimal solution) for various methods across 29 TSPLIB instances. For a fair comparison, we report results for EDISCO with 4x sampling decoding and 2-OPT post-processing, following the same setting as in (Li et al., 2023). Results for other baselines are from (Fu et al., 2021).

EDISCO achieves the lowest average optimality gap of 0.088%, representing a 31.6% relative improvement over the previous best method T2T (0.133%). Notably, EDISCO obtains optimal solutions (0.000% gap) on 6 instances and near-optimal solutions (< 0.05% gap) on 19 out of 29 instances. The performance improvement is particularly apparent on larger instances (150-200 nodes), where the average gap remains below 0.15%.

F.4 NOISE SCHEDULE DESIGN FOR TSP DIFFUSION

We conducted a comprehensive comparison of different noise schedule designs for the continuous-time categorical diffusion model applied to TSP. The noise schedule $\beta(t)$ controls the rate of information destruction during the forward diffusion process and significantly impacts model performance.

We evaluated three families of noise schedules:

Linear Schedule: Following the standard approach in diffusion models (Ho et al., 2020), we tested linear schedules with:

$$\beta(t) = \beta_{\min} + t(\beta_{\max} - \beta_{\min}) \quad (19)$$

1458
 1459 Table 7: Solution quality for methods trained on random 100-node problems and evaluated on
 1460 TSPLIB instances with 50-200 nodes. Results of DIFUSCO and T2T are from (Li et al., 2023),
 1461 which are based on 4x sampling decoding with 2-OPT post-processing. Results of other baselines
 1462 are from Fu et al. (2021). Bold indicates the best performance, and underlined indicates the second-best.
 1463

| Instance | AM | GCN | Learn2OPT | GNNGLS | DIFUSCO | T2T | EDISCO (Ours) |
|-----------------|-----------|------------|------------------|---------------|----------------|---------------|----------------------|
| eil51 | 16.767% | 40.025% | 1.725% | 1.529% | 0.314% | 0.314% | 0.217% |
| berlin52 | 4.169% | 33.225% | 0.449% | 0.142% | 0.000% | 0.000% | 0.000% |
| st70 | 1.737% | 24.785% | 0.040% | 0.764% | 0.172% | 0.000% | 0.000% |
| eil76 | 1.992% | 27.411% | 0.096% | 0.163% | 0.217% | <u>0.163%</u> | 0.108% |
| pr76 | 0.816% | 27.702% | 1.228% | 0.039% | 0.043% | <u>0.039%</u> | 0.024% |
| rat99 | 2.645% | 17.633% | 0.123% | 0.550% | 0.016% | 0.000% | 0.000% |
| kroA100 | 4.017% | 28.828% | 18.313% | 0.728% | 0.050% | 0.000% | 0.000% |
| kroB100 | 5.142% | 34.668% | 1.119% | 0.147% | 0.006% | 0.000% | 0.003% |
| kroC100 | 0.972% | 35.506% | 0.349% | 1.571% | 0.000% | 0.000% | 0.000% |
| kroD100 | 2.717% | 38.018% | 0.866% | 0.572% | 0.000% | 0.000% | 0.002% |
| kroE100 | 1.470% | 26.568% | 1.832% | 0.140% | 0.000% | 0.000% | 0.000% |
| rd100 | 3.407% | 50.432% | 1.725% | 0.003% | 0.000% | 0.000% | 0.001% |
| eil101 | 2.994% | 26.701% | 0.387% | 1.529% | 0.124% | 0.000% | 0.008% |
| lin105 | 1.739% | 34.902% | 1.867% | 0.484% | 0.441% | 0.393% | 0.267% |
| pr107 | 3.933% | 80.564% | 0.898% | 0.439% | 0.714% | 0.155% | 0.093% |
| pr124 | 3.677% | 70.146% | 10.322% | 0.755% | 0.997% | 0.584% | 0.372% |
| bier127 | 5.908% | 45.561% | 3.044% | 1.948% | 1.064% | 0.718% | 0.481% |
| ch130 | 3.182% | 39.090% | 0.709% | 3.519% | <u>0.077%</u> | <u>0.077%</u> | 0.046% |
| pr136 | 5.064% | 58.673% | 0.000% | 3.387% | 0.182% | 0.000% | 0.004% |
| pr144 | 7.641% | 55.837% | 1.526% | 3.581% | 1.816% | 0.000% | 0.011% |
| ch150 | 4.584% | 49.743% | 0.312% | 2.113% | 0.473% | 0.324% | 0.218% |
| kroA150 | 3.784% | 45.411% | 0.724% | 2.984% | <u>0.193%</u> | <u>0.193%</u> | 0.117% |
| kroB150 | 2.437% | 56.743% | 0.086% | 3.258% | 0.366% | 0.021% | 0.013% |
| pr152 | 7.494% | 33.925% | 0.029% | 3.119% | <u>0.687%</u> | <u>0.687%</u> | 0.428% |
| u159 | 7.551% | 63.338% | 10.534% | 1.020% | 0.000% | 0.000% | 0.003% |
| rat195 | 6.893% | 24.968% | 0.743% | 1.666% | 0.887% | 0.018% | 0.012% |
| d198 | 373.020% | 62.351% | 0.522% | 4.772% | 0.000% | 0.000% | 0.006% |
| kroA200 | 7.106% | 40.885% | 1.441% | 2.029% | 0.259% | 0.000% | 0.007% |
| kroB200 | 8.541% | 43.643% | 2.064% | 2.589% | <u>0.171%</u> | <u>0.171%</u> | 0.114% |
| Mean | 16.767% | 40.025% | 1.725% | 1.529% | 0.319% | 0.133% | 0.088% |

1490
 1491 **Exponential Schedule:** Based on Campbell et al. (Campbell et al., 2022), we evaluated exponential
 1492 schedules:
 1493

$$\beta(t) = ab^t \log(b) \quad (20)$$

1494
 1495 **Cosine Schedule:** Following Sun et al. (Sun et al., 2023b), we tested cosine schedules with im-
 1496 proved numerical stability:
 1497

$$\beta(t) = \text{clip}\left(\frac{\pi}{4} \cdot \frac{\tan(\pi t/2)}{\sqrt{\cos(\pi t/2) + \epsilon}}, \beta_{\min}, \beta_{\max}\right) \quad (21)$$

1500
 1501 We trained each schedule variant for 50 epochs on TSP-50 using 10,000 randomly generated in-
 1502 stances for fast verification. The same network architecture and hyperparameters were maintained
 1503 across all experiments. Each schedule family was tested with three different parameterizations:
 1504 baseline (standard parameters), conservative (slower noise injection with smaller β values), and
 1505 aggressive (faster noise injection with larger β values).

1506
 1507 Table 8 presents the comprehensive results. Linear schedules demonstrated the best overall perfor-
 1508 mance, with the baseline configuration ($\beta_{\min} = 0.1, \beta_{\max} = 1.5$) achieving the lowest validation
 1509 gap of 2.29%. The aggressive variant ($\beta_{\max} = 2.0$) and conservative variant ($\beta_{\max} = 1.0$) both un-
 1510 derperformed at 2.88% and 2.74% respectively, suggesting that moderate noise injection is optimal
 1511 for TSP diffusion.

Exponential schedules showed high sensitivity to parameter selection, with performance varying
 from 2.60% to 4.19% across configurations. Cosine schedules consistently underperformed with

| Schedule | Configuration | Optimality Gap (%) | | Conv. | Inference |
|-------------|---------------------------------------|--------------------|-------------|-------|-------------|
| Type | | Best | Final | Epoch | Time (s) |
| Linear | Aggressive: $\beta \in [0.1, 2.0]$ | 2.88 | 3.03 | 50 | 1.06 |
| | Baseline: $\beta \in [0.1, 1.5]$ | 2.29 | 2.63 | 45 | 1.06 |
| | Conservative: $\beta \in [0.1, 1.0]$ | 2.74 | 3.51 | 50 | 1.16 |
| Exponential | Baseline: $a=0.5, b=4.0$ | 3.39 | 2.60 | 50 | 1.04 |
| | Conservative: $a=0.3, b=3.0$ | 4.19 | 4.42 | 50 | 1.13 |
| | Aggressive: $a=0.8, b=5.0$ | 2.60 | 3.78 | 50 | 1.05 |
| Cosine | Aggressive: $\beta \in [0.001, 10.0]$ | 4.83 | 5.37 | 45 | 1.05 |
| | Baseline: $\beta \in [0.01, 5.0]$ | 3.45 | 4.51 | 40 | 1.07 |
| | Conservative: $\beta \in [0.1, 3.0]$ | 3.25 | 4.09 | 40 | 1.06 |

gaps ranging from 3.25% to 4.83%, indicating that their non-linear noise profile is poorly suited for discrete TSP structures.

The superiority of linear schedules in TSP contrasts with image generation, where cosine schedules often excel (Nichol & Dhariwal, 2021). We attribute this to the discrete nature of TSP adjacency matrices, which benefit from uniform, predictable noise injection rather than variable rates. These findings validate our choice of $\beta_{\min} = 0.1, \beta_{\max} = 1.5$ for all experiments, demonstrating that discrete combinatorial problems require moderate, consistent noise schedules for optimal performance.

F.5 EVALUATION ON ADAPTIVE MIXING PARAMETERS

We conduct a comprehensive evaluation on the adaptive mixing strategy parameters to justify our design choices. The adaptive mixing strategy (Equation 7) balances between diffusion-based transitions and direct model predictions using a time-dependent weight function $w(t)$, with deterministic switching near $t = 0$.

Mixing Weight Functions We evaluate different weight functions $w(t)$ that control the interpolation between diffusion dynamics and predicted \mathbf{X}_0 :

Table 9: Comparison of mixing weight functions on TSP-50 and TSP-100. All models use 50 diffusion steps with greedy decoding.

| Weight Function | TSP-50 | | TSP-100 | | Convergence Epoch |
|-----------------------------------|-------------|----------|-------------|----------|-------------------|
| | Gap (%) ↓ | Time (s) | Gap (%) ↓ | Time (s) | |
| Linear: $w(t) = t$ | 0.01 | 1.06 | 0.04 | 2.84 | 35 |
| Quadratic: $w(t) = t^2$ | 0.03 | 1.08 | 0.08 | 2.87 | 38 |
| Square root: $w(t) = \sqrt{t}$ | 0.02 | 1.07 | 0.06 | 2.85 | 36 |
| Cosine: $w(t) = \cos(\pi t/2)$ | 0.04 | 1.09 | 0.09 | 2.89 | 40 |
| Exponential: $w(t) = e^{-2(1-t)}$ | 0.05 | 1.11 | 0.11 | 2.91 | 42 |
| Constant: $w(t) = 0.5$ | 0.18 | 1.05 | 0.42 | 2.82 | 48 |
| No mixing ($w(t) = 1$) | 0.31 | 1.04 | 0.68 | 2.81 | 52 |
| Pure prediction ($w(t) = 0$) | 0.28 | 1.04 | 0.46 | 2.81 | 50 |

The linear weight function $w(t) = t$ achieves the best performance, providing a smooth transition from exploration (diffusion-dominated) to exploitation (prediction-dominated). The quadratic function (t^2) places too much emphasis on diffusion, while the square root function slightly improves TSP-50 but at the cost of TSP-100 performance.

Deterministic Switching Thresholds We evaluate different thresholds for switching to deterministic argmax selection:

1566

Table 10: Effect of deterministic switching thresholds on solution quality. Models use linear mixing $w(t) = t$. \dagger Percentage of instances where the greedy decoder fails to construct a valid Hamiltonian cycle.

| Time Threshold | Step Threshold | TSP-50 | TSP-100 | TSP-500 | Failed Tours \dagger |
|----------------|---------------------|-------------|-------------|-------------|------------------------|
| $t < 0.05$ | $ \Delta t < 0.01$ | 0.04 | 0.09 | 2.41 | 3.2% |
| $t < 0.1$ | $ \Delta t < 0.02$ | 0.01 | 0.04 | 1.95 | 0.0% |
| $t < 0.15$ | $ \Delta t < 0.03$ | 0.02 | 0.05 | 1.98 | 0.0% |
| $t < 0.2$ | $ \Delta t < 0.04$ | 0.03 | 0.07 | 2.12 | 0.0% |
| $t < 0.25$ | $ \Delta t < 0.05$ | 0.06 | 0.13 | 2.38 | 0.1% |
| No switching | No switching | 0.08 | 0.21 | 2.94 | 1.8% |

1577

1578

The threshold $t < 0.1$ with $|\Delta t| < 0.02$ provides the optimal balance. Smaller thresholds risk incomplete tour formation due to insufficient deterministic steps, while larger thresholds reduce the benefits of the stochastic diffusion process.

1582

Joint Impact Analysis We evaluate the joint effect of mixing function and switching threshold on TSP-500:

1585

1586

Table 11: Joint ablation of mixing function and deterministic threshold on TSP-500 (optimality gap %).

| Mixing Function | Deterministic Threshold | | | |
|-------------------|-------------------------|-------------|------------|-----------|
| | $t < 0.05$ | $t < 0.1$ | $t < 0.15$ | $t < 0.2$ |
| $w(t) = t$ | 2.41 | 1.95 | 1.98 | 2.12 |
| $w(t) = t^2$ | 2.68 | 2.23 | 2.19 | 2.25 |
| $w(t) = \sqrt{t}$ | 2.33 | 2.01 | 2.04 | 2.18 |
| $w(t) = 0.5$ | 3.12 | 2.86 | 2.91 | 3.05 |

1595

1596

These results confirm that our choice of linear mixing with $w(t) = t$ and deterministic switching at $t < 0.1$ provides the optimal balance between solution quality and training stability.

1599

F.6 EVALUATION ON ARCHITECTURAL HYPERPARAMETERS

1602

1603

We conduct systematic evaluations of critical architectural hyperparameters to justify our design choices.

1604

1605

Step Size α for Coordinate Updates The step size α in Equation 9 controls the magnitude of coordinate updates during message passing. We evaluate different values on TSP-50:

1607

1608

Table 12: Effect of step size α on model performance and training stability (TSP-50). *Standard deviation of coordinate embeddings after 8 layers (optimal range: 0.3-0.4).

| Step Size α | 0.01 | 0.05 | 0.1 | 0.2 | 0.5 |
|--------------------|------|------|-------------|----------|-----------|
| Optimality Gap (%) | 0.08 | 0.03 | 0.01 | 0.15 | Diverged |
| Coordinate Std* | 0.42 | 0.38 | 0.35 | 0.18 | 0.02 |
| Training Stable | ✓ | ✓ | ✓ | Unstable | Collapsed |
| Convergence Epoch | 42 | 38 | 35 | 48 | - |

1616

1617

1618

1619

As shown in Table 12, smaller α values (0.01, 0.05) maintain stability but converge more slowly and achieve suboptimal performance. Larger values ($\alpha \geq 0.2$) cause coordinate collapse, where the standard deviation of coordinate embeddings approaches zero, indicating all cities converge to similar positions in the latent space.

1620 **Temperature Parameter τ for Weight Scaling** The temperature parameter τ in Equation 9 scales
 1621 the MLP outputs before applying tanh, preventing gradient saturation:
 1622
 1623

1624 Table 13: Effect of temperature τ on gradient flow and performance (TSP-50). † Average gradient
 1625 norm in coordinate MLP during first 10 epochs. ‡ Percentage of tanh outputs with $|w_{ij}| > 0.95$.

| Temperature τ | 1 | 5 | 10 | 20 | 50 |
|----------------------------------|-------|-------|--------------|-------|-------|
| Optimality Gap (%) | 0.12 | 0.04 | 0.01 | 0.02 | 0.05 |
| Avg. Gradient Norm † | 0.003 | 0.018 | 0.042 | 0.038 | 0.031 |
| Tanh Saturation Rate ‡ | 68% | 24% | 8% | 12% | 18% |
| Convergence Epoch | 52 | 40 | 35 | 36 | 39 |

1626 Table 13 shows that $\tau = 10$ achieves the optimal balance. Lower values ($\tau \leq 5$) cause excessive
 1627 saturation, leading to vanishing gradients and slower convergence. Higher values ($\tau \geq 20$) reduce
 1628 the non-linearity's effectiveness, diminishing the model's expressiveness.
 1629
 1630

1631 **Joint Impact Analysis** We evaluate the joint effect of α and τ on TSP-100 performance:

1632 Table 14: Joint ablation of α and τ on TSP-100 (optimality gap %).

| α | τ | 1 | 5 | 10 | 20 | 50 |
|----------|-----------|-----------|-------------|----------|----------|----|
| 0.01 | 0.21 | 0.15 | 0.08 | 0.09 | 0.12 | |
| 0.05 | 0.18 | 0.09 | 0.05 | 0.06 | 0.08 | |
| 0.1 | 0.15 | 0.06 | 0.04 | 0.05 | 0.07 | |
| 0.2 | Unstable | 0.18 | 0.15 | 0.16 | 0.19 | |
| 0.5 | Collapsed | Collapsed | Diverged | Diverged | Diverged | |

1640 The joint ablation confirms that $\alpha = 0.1$ and $\tau = 10$ provide the optimal configuration, with
 1641 performance degrading smoothly as we deviate from these values.
 1642
 1643

1644 F.7 MODEL EFFICIENCY

1645 Although the results in the main text are from the full-scale EDISCO model, it is interesting to
 1646 see EDISCO's performance under reduced model sizes. Table 15 presents a comprehensive com-
 1647 parison of model efficiency across different architectures and scales, demonstrating how EDISCO's
 1648 equivariant design enables strong performance even with reduced model capacity and training data.
 1649
 1650

1651 Table 15: Model efficiency and performance comparison across TSP scales. Training instances
 1652 shown in thousands (K) with corresponding optimality gaps (%). EDISCO-Full uses 12 EGNN
 1653 layers with 256 hidden dimension (5.5M parameters), EDISCO-Medium uses 12 layers with 128
 1654 hidden dimension (2.6M parameters), and EDISCO-Small uses 8 layers with 128 hidden dimension
 1655 (1.4M parameters). † 12-layer GNN with width 256. * Same architecture as DIFUSCO.

| Method | Model Parameters | Training Instances (K) / Optimality Gap (%) | | | | |
|--------------------|------------------|---|-------------------|------------------|------------------|-----------------|
| | | TSP-50 | TSP-100 | TSP-500 | TSP-1000 | TSP-10000 |
| DIFUSCO † | 5.3M | 1500 / 0.48 | 1500 / 1.01 | 128 / 9.41 | 64 / 11.24 | 6.4 / 8.95 |
| DISCO † | 5.3M | 1500 / 0.16 | 1500 / 0.58 | - / - | - / - | 6.4 / 2.90 |
| T2T * | 5.3M | 1500 / 0.04 | 1500 / 0.18 | 128 / 5.09 | 64 / 8.87 | 6.4 / 2.92 |
| EDISCO-Full | 5.5M | 500 / 0.01 | 500 / 0.04 | 60 / 1.95 | 30 / 2.85 | 3 / 1.98 |
| EDISCO-Medium | 2.6M | 500 / 0.03 | 500 / 0.08 | 60 / 2.18 | 30 / 2.85 | 3 / 2.43 |
| EDISCO-Small | 1.4M | 300 / 0.08 | 300 / 0.7 | 40 / 4.18 | 20 / 5.21 | 2 / 3.18 |

1672 EDISCO-Full (5.5M parameters) uses 3 \times less training data than baselines on small instances (500K
 1673 vs 1.5M) and 2 \times less on large instances. It achieves 0.01% gap on TSP-50 compared to DIFUSCO's

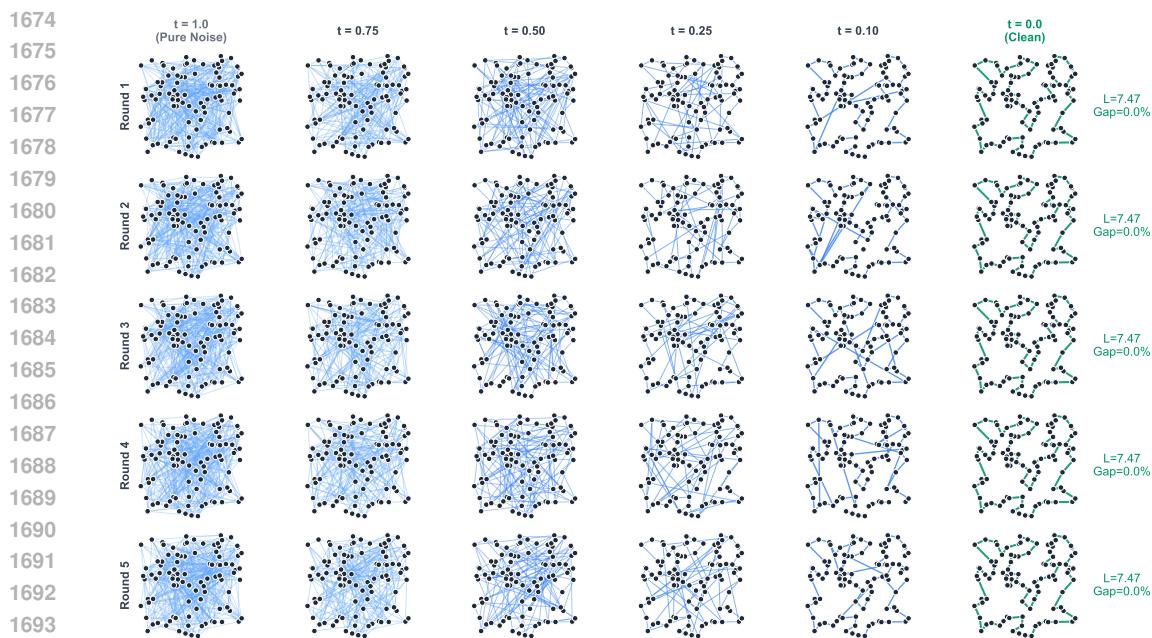


Figure 4: Five rounds of the denoising process of EDISCO on the original TSP instance.

0.48%, and 1.95% on TSP-500 versus DIFUSCO’s 9.41%. On TSP-10000, EDISCO-Full achieves 1.98% gap with 3K training instances, outperforming DIFUSCO’s 8.95% gap with 6.4K instances.

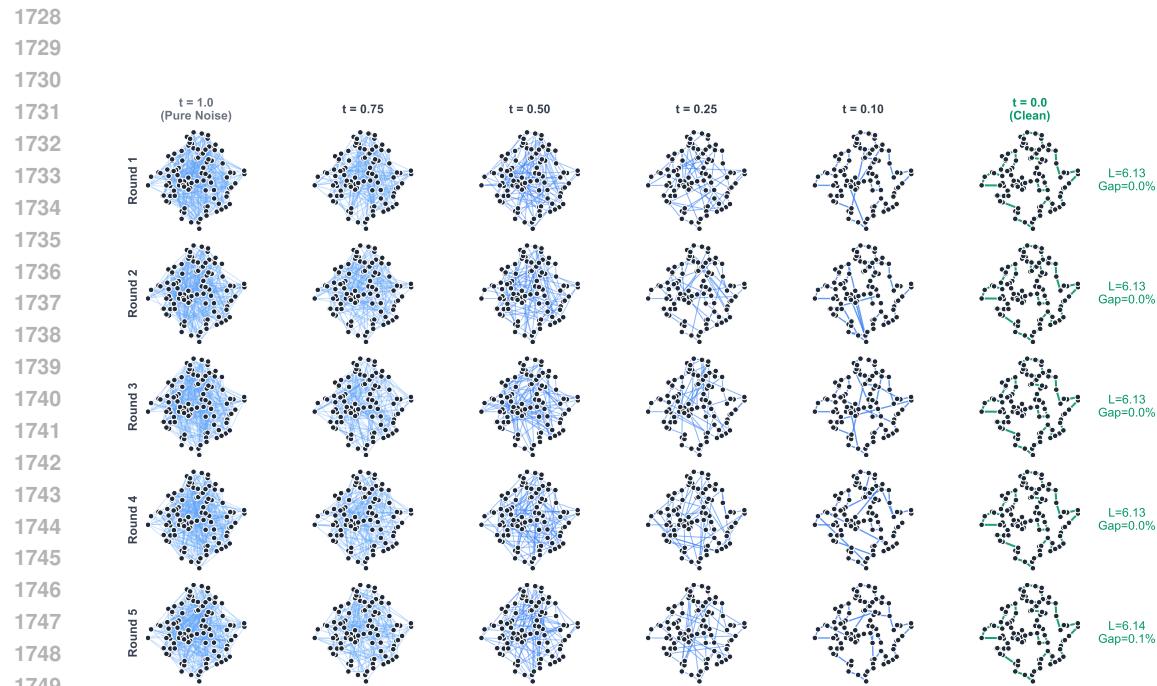
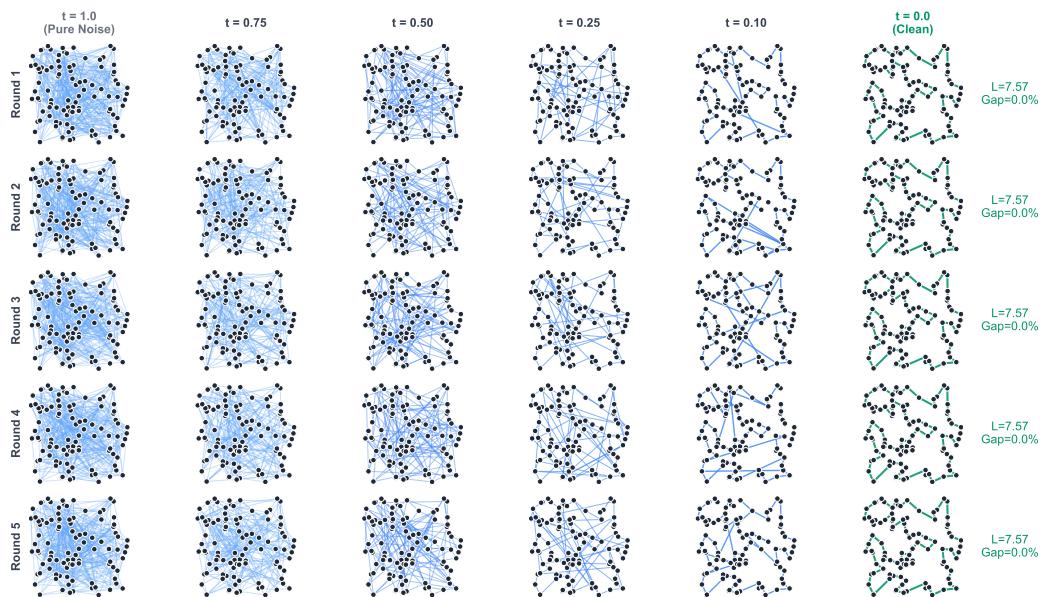
EDISCO-Medium (2.6M parameters), with less than half the parameters of baseline models, achieves 0.03% gap on TSP-50 and 2.18% on TSP-500. It matches EDISCO-Full’s performance on TSP-1000 at 2.85% gap. Compared to T2T, which achieves 0.04% on TSP-50 with 5.3M parameters, EDISCO-Medium achieves comparable performance (0.03%) with 2.6M parameters and the same 500K training instances.

EDISCO-Small (1.4M parameters) uses 3.8x fewer parameters than baselines and requires the least training data: 300K for TSP-50/100, 40K for TSP-500, and 2K for TSP-10000. It achieves 0.08% gap on TSP-50 and 0.7% on TSP-100. On TSP-500, with 40K training instances, it achieves 4.18% gap, compared to DIFUSCO’s 9.41% with 128K instances. On TSP-10000, EDISCO-Small achieves 3.18% gap, outperforming DIFUSCO (8.95%) and DISCO (2.90%).

The results for the amount of training data represent the minimum numbers required to ensure the EDISCO converges to the optimal gaps. After this, even with increased training data, there is no noticeable improvement in the optimality gaps.

F.8 VISUALIZATION OF EDISCO’s E(2)-EQUIVARIANCE

Figure 4 illustrates the denoising process of EDISCO on a standard TSP-100 instance. To empirically validate the E(2) equivariance of our architecture, we present the same denoising process on rotated versions of the instance in Figures 5 and 6. Only the greedy decoder is used without any other post-processing techniques. The visualization shows progression from pure noise ($t = 1.0$) to clean tours ($t = 0.0$) across five independent sampling rounds, demonstrating consistent convergence to high-quality solutions. The similar performance across all three orientations confirms that EDISCO has learned truly rotation-invariant representations.

Figure 5: Five rounds of the denoising process of EDISCO on the 45° rotated instance.Figure 6: Five rounds of the denoising process of EDISCO on the 90° rotated instance.